

Adobe ACE Flex 3 & Air Study Guide

By Jonnie Spratley 1/9/09

Here is all of the tested content that is covered on the ACE Flex 3 & Air Exam.

I have done all of the research for every section and topic that is said to be on the exam. Now it is up to you to study this guide and pass the test to become a certified Flex 3 & Air Developer.

Exam Structure

The following lists the topic areas and percentage of questions delivered in each topic area:

Topic Area	% Of Exam	# Of Questions
Creating a User Interface (UI)	22	11
Flex system architecture and design	18	9
Programming Flex applications with ActionScript	24	12
Interacting with data sources and servers	16	8
Using Flex in the Adobe Integrated Runtime (AIR)	20	10

Number of Questions and Passing Score

- 50 questions
- 67% minimum required to pass

Test Content: Topic Areas and Objectives

Following is a detailed outline of the information covered on the exam.

1. Creating a User Interface (UI)
 - 1.1 - Identify and describe the basic UI controls used in a Flex application.
 - 1.2 - Identify the purpose of UI containers and when to use them.
 - 1.3 - Change the look and feel of a design by using API styles, style sheets, filters, and blends.
 - 1.4 - Dynamically change the look of an application by using view states, transitions and effects.
 - 1.5 - Position UI elements by using constraint-based layout.
 - 1.6 - Position UI elements by using enhanced constraints.
 - 1.7 - Implement application navigation by using navigator containers.
 - 1.8 - Customize list-based controls.

2. Flex system architecture and design
 - 2.1 - Create and use custom components.
 - 2.2 - Transfer data between components by using data bindings.
 - 2.3 - Create, handle, and dispatch custom events.
 - 2.4 - Handle framework events.
 - 2.5 - List and describe the differences between model, view, and controller code in a Flex application.
3. Programming Flex applications with ActionScript
 - 3.1 - Define and extend an ActionScript class.
 - 3.2 - Implement an ActionScript interface.
 - 3.3 - Use access modifiers with classes and class members.
 - 3.4 - Under the purpose of and implement data transfer objects.
 - 3.5 - Implement accessor methods in ActionScript.
 - 3.6 - Use an ArrayCollection to sort, filter, and provide data.
 - 3.7 - Implement data validation.
 - 3.8 - Manipulate XML data by using E4X.
4. Interacting with data sources and servers
 - 4.1 - Implement simple LiveCycle Data Services (LCDS) messaging and data management.
 - 4.2 - Create, connect to, and define a local database.
 - 4.3 - Add, update, and remove records from local database.
 - 4.4 - Interact with remote data and services by using Remote Procedure Call (RPC) services.
 - 4.5 - Upload files to a server.
5. Using Flex in the Adobe Integrated Runtime (AIR)
 - 5.1 - Given a scenario, compile and export a release build of an AIR application.
 - 5.2 - Create, populate, and delete files and directories on a local file system.
 - 5.3 - Create and customize native windows and menus.
 - 5.4 - Adding drag-and-drop functionality to and from the desktop.
 - 5.5 - Install, uninstall, and update an AIR application.
 - 5.6 - List and describe the AIR security contexts.

The following is a detailed outline of the information covered on the exam with all of the answers in full detail.

1. Creating a User Interface (UI)

1.1 - Identify and describe the basic UI controls used in a Flex application.

a. All Basic UI Controls

Component	Description
-----------	-------------

<p>Button</p>	<p>Buttons typically use event listeners to perform an action when the user selects the control.</p> <p>When a user clicks the mouse on a Button control, and the Button control is enabled, it dispatches a click event and a <code>buttonDown</code> event.</p> <p>A button always dispatches events such as the <code>mouseMove</code>, <code>mouseOver</code>, <code>mouseOut</code>, <code>rollOver</code>, <code>rollOut</code>, <code>mouseDown</code>, and <code>mouseUp</code> events whether enabled or disabled.</p>
<p>PopUpButton</p>	<p>The PopUpButton control consists of two horizontal buttons: a main button, and a smaller button called the pop-up button, which only has an icon.</p> <p>The main button is a Button control. The pop-up button, when clicked, opens a second control called the pop-up control.</p> <p>Clicking anywhere outside the PopUpButton control, or in the pop-up control, closes the pop-up control</p>
<p>ButtonBar and ToggleButtonBar</p>	<p>The ButtonBar and ToggleButtonBar controls define a horizontal or vertical row of related buttons with a common appearance. The controls define a single event, the <code>itemClick</code> event, that is dispatched when any button in the control is selected.</p> <p>The ButtonBar control defines group of buttons that do not retain a selected state. When you select a button in a ButtonBar control, the button changes its appearance to the selected state; when you release the button, it returns to the deselected state.</p> <p>The ToggleButtonBar control defines a group buttons that maintain their state, either selected or deselected. Only one button in the ToggleButtonBar control can be in the selected state.</p> <p>That means when you select a button in a ToggleButtonBar control, the button stays in the selected state until you select a different button.</p>

LinkBar	<p>A LinkBar control defines a horizontal or vertical row of LinkButton controls that designate a series of link destinations.</p> <p>You typically use a LinkBar control to control the active child container of a ViewStack container, or to create a standalone set of links.</p>
TabBar	<p>As with the LinkBar control, you can use a TabBar control to control the active child container of a ViewStack container. The syntax for using a TabBar control to control the active child of a ViewStack container is the same as for a LinkBar control.</p>
CheckBox	<p>The CheckBox control is a commonly used graphical control that can contain a check mark or be unchecked (empty). You can use CheckBox controls wherever you need to gather a set of true or false values that aren't mutually exclusive. You can add a text label to a CheckBox control and place it to the left, right, top, or bottom. Flex clips the label of a CheckBox control to fit the boundaries of the control.</p>
RadioButton	<p>The RadioButton control is a single choice in a set of mutually exclusive choices.</p> <p>A RadioButton group is composed of two or more RadioButton controls with the same group name. Only one member of the group can be selected at any given time. Selecting an unselected group member deselects the currently selected RadioButton control in the group.</p>
NumericStepper	<p>You can use the NumericStepper control to select a number from an ordered set.</p> <p>The NumericStepper control consists of a single-line input text field and a pair of arrow buttons for stepping through the valid values; you can also use the Up Arrow and Down Arrow keys to cycle through the values.</p>
DateChoose and DateField	<p>The DateChooser and DateField controls let users select dates from graphical calendars. The DateChooser control user interface is the calendar.</p> <p>The DateField control has a text field that uses a date chooser popup to select the date as a result. The DateField properties are a superset of the DateChooser properties.</p>

LinkButton	<p>The LinkButton control creates a single-line hypertext link that supports an optional icon. You can use a LinkButton control to open a URL in a web browser.</p>
HSlider and VSlider	<p>You can use the slider controls to select a value by moving a slider thumb between the end points of the slider track.</p> <p>The current value of the slider is determined by the relative location of the thumb between the end points of the slider, corresponding to the slider's minimum and maximum values. By default, the minimum value of a slider is 0 and the maximum value is 10.</p> <p>The current value of the slider can be any value in a continuous range between the minimum and maximum values, or it can be one of a set of discrete values, depending on how you configure the control.</p>
SWFLoader	<p>The SWFLoader control lets you load one Flex application into another Flex application as a SWF file. It has properties that let you scale its contents. It can also resize itself to fit the size of its contents. By default, content is scaled to fit the size of the SWFLoader control.</p> <p>The SWFLoader control can also load content on demand programmatically, and monitor the progress of a load operation.</p> <p>The SWFLoader control also lets you load the contents of a GIF, JPEG, PNG, SVG, or SWF file into your application, where the SWF file does not contain a Flex application, or a ByteArray representing a SWF, GIF, JPEG, or PNG.</p>
Image	<p>Adobe Flex supports several image formats, including GIF, JPEG, PNG, SVG, and SWF files.</p> <p>You can import these images into your applications by using the Image control.</p>
VideoDisplay	<p>Flex supports the VideoDisplay control to incorporate streaming media into Flex applications.</p> <p>Flex supports the Flash Video File (FLV) file format with this control.</p>

ColorPicker	<p>The ColorPicker control lets users select a color from a drop-down swatch panel (palette). It initially appears as a preview sample with the selected color.</p> <p>When a user selects the control, a color swatch panel appears. The panel includes a sample of the selected color and a color swatch panel.</p> <p>By default, the swatch panel displays the web-safe colors (216 colors, where each of the three primary colors has a value that is a multiple of 33, such as #CC0066).</p>
Alert	<p>All Flex components can call the static <code>show()</code> method of the Alert class to open a pop-up modal dialog box with a message and an optional title, buttons, and icons.</p>
ProgressBar	<p>The ProgressBar control provides a visual representation of the progress of a task over time.</p> <p>There are two types of ProgressBar controls:</p> <p>Determinate and Indeterminate.</p> <p>A determinate ProgressBar control is a linear representation of the progress of a task over time. You can use this when the user is required to wait for an extended period of time, and the scope of the task is known.</p> <p>An indeterminate ProgressBar control represents time-based processes for which the scope is not yet known. As soon as you can determine the scope, you should use a determinate ProgressBar control.</p>
HRule and VRule	<p>The HRule (Horizontal Rule) control creates a single horizontal line and the VRule (Vertical Rule) control creates a single vertical line.</p> <p>You typically use these controls to create dividing lines within a container.</p>
Scrollbar	<p>The VScrollBar (vertical ScrollBar) control and HScrollBar (horizontal ScrollBar) controls let the user control the portion of data that is displayed when there is too much data to fit in the display area.</p>

1.2 - Identify the purpose of UI containers and when to use them.

a. All UI Containers And There Purpose

Container	Purpose
Canvas	<p>A Canvas layout container defines a rectangular region in which you place child containers and controls. It is the only container that lets you explicitly specify the location of its children within the container by using the x and y properties of each child.</p> <p>Flex sets the children of a Canvas layout container to their preferred width and preferred height. You may override the value for a child's preferred width by setting its width property to either a fixed pixel value or a percentage of the container size. You can set the preferred height in a similar manner.</p> <p>If you use percentage sizing inside a Canvas container, some of your components may overlap. If this is not the effect you want, plan your component positions and sizes carefully.</p> <p>The Canvas container has the following default sizing characteristics:</p> <p>Default size: Large enough to hold all its children at the default sizes of the children</p> <p>Default padding: 0 pixels for the top, bottom, left, and right values</p>

<p>Box, HBox, VBox</p>	<p>A Box container lays out its children in a single vertical column or a single horizontal row. The direction property determines whether to use vertical (default) or horizontal layout.</p> <p>The Box class is the base class for the VBox and HBox classes. You use the <code><mx:Box></code>, <code><mx:VBox></code>, and <code><mx:HBox></code> tags to define Box containers.</p> <p>A Box container has the following default sizing characteristics:</p> <p>Default size:</p> <p>Vertical Box: The height is large enough to hold all its children at the default or explicit height of the children, plus any vertical gap between the children, plus the top and bottom padding of the container. The width is the default or explicit width of the widest child, plus the left and right padding of the container.</p> <p>Horizontal Box: The width is large enough to hold all of its children at the default width of the children, plus any horizontal gap between the children, plus the left and right padding of the container. The height is the default or explicit height of the tallest child, plus the top and bottom padding for the container.</p> <p>Default padding: 0 pixels for the top, bottom, left, and right values.</p>
	<p>The ControlBar container lets you place controls at the bottom of a Panel or TitleWindow container. The <code><mx:ControlBar></code> tag must be the last child tag of the enclosing tag for the Panel or TitleWindow container.</p> <p>The ControlBar is a Box with a background and default style properties.</p> <p>A ControlBar container has the following default sizing characteristics:</p> <p>Default size: Height is the default or explicit height of the tallest child, plus the top and bottom padding of the container. Width is large enough to hold all of its children at the default or explicit width of the children, plus any horizontal gap between the children, plus the left and right padding of the container.</p> <p>Default padding: 10 pixels for the top, bottom, left, and right values.</p>

ApplicationControlBar	<p>You use the ApplicationControlBar container to hold components that provide access to application navigation elements and commands.</p> <p>The ApplicationControlBar container can be in either of the following modes:</p> <p>Docked mode: The bar is always at the top of the application's drawing area and becomes part of the application chrome. Any application-level scroll bars don't apply to the component, so that it always remains at the top of the visible area, and the bar expands to fill the width of the application. To create a docked bar, set the value of the dock property to <code>true</code>.</p> <p>Normal mode: The bar can be placed anywhere in the application, gets sized and positioned just like any other component, and scrolls with the application. To create a normal bar, set the value of the dock property to <code>false</code> (default).</p>
DividedBox, HDividedBox, VDividedBox	<p>A DividedBox container measures and lays out its children horizontally or vertically in exactly the same way as a Box container, but it inserts draggable dividers in the gaps between the children. Users can drag any divider to resize the children on each side.</p> <p>The DividedBox class is the base class for the more commonly used HDividedBox and VDividedBox classes.</p> <p>The direction property of a DividedBox container, inherited from Box container, determines whether it has horizontal or vertical layout.</p> <p>A DividedBox, HDividedBox, or VDividedBox container has the following default sizing characteristics:</p> <p>Default size:</p> <p>Vertical DividedBox: Height is large enough to hold all of its children at the default or explicit heights of the children, plus any vertical gap between the children, plus the top and bottom padding of the container. The width is the default or explicit width of the widest child, plus the left and right padding of the container.</p> <p>Horizontal DividedBox: Width is large enough to hold all of its children at the default or explicit widths of the children, plus any horizontal gap between the children, plus the left and right padding of the container. Height is the default or explicit height of the tallest child plus the top and bottom padding of the container.</p> <p>Default padding: 0 pixels for the top, bottom, left, and right values.</p> <p>Default gap: 10 pixels for the horizontal and vertical gaps.</p>

Form, FormHeading, FormItem	<p>The Form container lets you control the layout of a form, mark form fields as required or optional, handle error messages, and bind your form data to the Flex data model to perform data checking and validation. It also lets you use style sheets to configure the appearance of your forms.</p> <p>The following table describes the components you use to create forms in Flex:</p> <p>Form: <code><mx:Form></code></p> <p>Defines the container for the entire form, including the overall form layout. Use the FormHeading control and FormItem container to define content. You can also insert other types of components in a Form container.</p> <p>FormHeading: <code><mx:FormHeading></code></p> <p>Defines a heading within your form. You can have multiple FormHeading controls within a single Form container.</p> <p>FormItem: <code><mx:FormItem></code></p> <p>Contains one or more form children arranged horizontally or vertically. Children can be controls or other containers. A single Form container can hold multiple FormItem containers.</p>
-----------------------------	--

Grid	<p>A Grid container lets you arrange children as rows and columns of cells, similar to an HTML table. The Grid container contains one or more rows, and each row can contain one or more cells, or items.</p> <p>You use the following tags to define a Grid control:</p> <p>The <code><mx:Grid></code> tag defines a Grid container.</p> <p>The <code><mx:GridRow></code> tag defines a grid row, which has one or more cells. The grid row must be a child of the <code><Grid></code> tag.</p> <p>The <code><mx:GridItem></code> tag defines a grid cell, and must be a child of the <code><GridRow></code> tag. The <code><mx:GridItem></code> tag can contain any number of children.</p> <p>The height of all the cells in a single row is the same, but each row can have a different height. The width of all cells in a single column is the same, but each column can have a different width. You can define a different number of cells for each row or each column of the Grid container. In addition, a cell can span multiple columns or multiple rows of the container.</p> <p>The Grid, GridRow, and GridItem containers have the following default sizing characteristics:</p> <p>Grid height: The sum of the default or explicit heights of all rows plus the gaps between rows.</p> <p>Grid width: The sum of the default or explicit width of all columns plus the gaps between columns.</p> <p>Height of each row and each cell: The default or explicit height of the tallest item in the row. If a GridItem container does not have an explicit size, its default height is the default or explicit height of the child in the cell.</p> <p>Width of each column and each cell: The default or explicit width of the widest item in the column. If a GridItem container does not have an explicit width, its default width is the default or explicit width of the child in the cell.</p> <p>Gap between rows and columns: Determined by the <code>horizontalGap</code> and <code>verticalGap</code> properties of the Grid class. The default value for both gaps is 6 pixels.</p> <p>Default padding: 0 pixels for the top, bottom, left, and right values, for all three container classes.</p>
------	--

Panel	<p>A Panel container consists of a title bar, a caption, a border, and a content area for its children. Typically, you use Panel containers to wrap top-level application modules. For example, you could include a shopping cart in a Panel container.</p> <p>The Panel container has the following default sizing characteristics:</p> <p>Default size: Height is large enough to hold all of its children at the default height of the children, plus any vertical gaps between the children, the top and bottom padding, the top and bottom borders, and the title bar. Width is the larger of the default width of the widest child plus the left and right padding of the container, or the width of the title text, plus the border.</p> <p>Padding: 4 pixels for the top, bottom, left, and right values.</p>
Tile	<p>A Tile container lays out its children in a grid of equal-sized cells. You can specify the size of the cells by using the <code>tileWidth</code> and <code>tileHeight</code> properties, or let the Tile container determine the cell size based on the largest child. A Tile container's <code>direction</code> property determines whether its cells are laid out horizontally or vertically, beginning from the upper-left corner of the Tile container.</p> <p>A Tile container has the following default sizing characteristics:</p> <p>Direction: horizontal</p> <p>Default size of all cells: Height is the default or explicit height of the tallest child. Width is the default or explicit width of the widest child. All cells have the same default size.</p> <p>Default size of Tile container: Flex computes the square root of the number of children, and rounds up to the nearest integer. For example, if there are 26 children, the square root is 5.1, which is rounded up to 6. Flex then lays out the Tile container in a 6 by 6 grid. Default height of the Tile container is equal to (tile cell default height) x (rounded square root of the number of children), plus any gaps between children and any padding. Default width is equal to (tile cell default width) x (rounded square root of the number of children), plus any gaps between children and any padding.</p> <p>Minimum size of Tile container: The default size of a single cell. Flex always allocates enough space to display at least one cell.</p> <p>Default padding: 0 pixels for the top, bottom, left, and right values.</p>

TitleWindow	<p>A TitleWindow layout container contains a title bar, a caption, a border, and a content area for its child. Typically, you use TitleWindow containers to wrap self-contained application modules. For example, you could include a form in a TitleWindow container. When the user completes the form, you can close the TitleWindow container programmatically, or let the user close it by using the Close button.</p> <p>The TitleWindow container has the following default sizing characteristics:</p> <p>Default size: Height is large enough to hold all of the children in the content area at the default or explicit heights of the children, plus the title bar and border, plus any vertical gap between the children, plus the top and bottom padding of the container. Width is the larger of the default or explicit width of the widest child, plus the left and right container borders padding, or the width of the title text.</p> <p>Borders: 10 pixels for the left and right values. 2 pixels for the top value. 0 pixels for the bottom value.</p> <p>Padding: 4 pixels for the top, bottom, left, and right values.</p>
-------------	--

1.3 - Change the look and feel of a design by using API styles, style sheets, filters, & blends.

a. Local style definitions

Use the `<mx:Style>` tag to define styles that apply to the current document and its children. You define styles in the `<mx:Style>` tag using CSS syntax and can define styles that apply to all instances of a control or to individual controls.

The following example defines a new style and applies it to only the `myButton` control

```
<mx:Style>
    .myFontStyle
    {
        fontSize: 15;
        color: #9933FF;
    }
</mx:Style>

<!-- This button has the custom style applied to it. -->
<mx:Button id="myButton" styleName="myFontStyle" label="Click Me"/>

<!-- This button uses default styles. -->
<mx:Button id="myButton2" label="Click Me"/>
```

The following example defines a new style that applies to all instances of the Button class:

```
<mx:Style>
    Button
    {
        fontSize: 15;
        color: #9933FF;
    }
</mx:Style>

<mx:Button id="myButton" label="Click Me"/>
<mx:Button id="myButton2" label="Click Me"/>
```

b. StyleManager class

Use the StyleManager class to apply styles to all classes or all instances of specified classes.

The following example sets the `fontSize` style to 15 and the `color` to 0x9933FF on all Button controls:

```
public function initApp():void
{
    StyleManager.getStyleDeclaration("Button").setStyle("fontSize",15);
    StyleManager.getStyleDeclaration("Button").setStyle("color",0x9933FF);
}
```

You can also use the `CSSStyleDeclaration` object to build run-time style sheets, and then apply them with the `StyleManager`'s `setStyleDeclaration()` method.

c. getStyle() and setStyle() methods

Use the `setStyle()` and `getStyle()` methods to manipulate style properties on instances of controls. Using these methods to apply styles requires a greater amount of processing power on the client than using style sheets but provides more granular control over how styles are applied.

The following example sets the `fontSize` to 15 and the `color` to 0x9933FF on only the `myButton` instance:

```
public function initApp():void
{
    myButton.setStyle("fontSize",15);
    myButton.setStyle("color",0x9933FF);
}
```

d. Inline styles

Use attributes of MXML tags to apply style properties. These properties apply only to the instance of the control. This is the most efficient method of applying instance properties because no ActionScript code blocks or method calls are required.

The following example sets the `fontSize` to 15 and the `color` to 0x9933FF on the `myButton` instance:

```
<!-- This button uses custom inline styles. -->
<mx:Button id="myButton" color="0x9933FF" fontSize="15" label="Click Me"/>

<!-- This button uses default styles. -->
<mx:Button id="myOtherButton" label="Click Me"/>
```

In an MXML tag, you must use the camel-case version of the style property. For example, you must use `fontSize` rather than `font-size` (the CSS convention) in the previous example.

v. Setting global styles

Most text and color styles, such as `fontSize` and `color`, are inheritable. When you apply an inheritable style to a container, all the children of that container inherit the value of that style property. If you set the color of a Panel container to green, all buttons in the Panel container are also green, unless those buttons override that color. Many styles, however, are not inheritable. If you apply them to a parent container, only the container uses that style. Children of that container do not use the values of non-inheritable styles. By using global styles, you can apply non-inheritable styles to all controls that do not explicitly override that style.

Flex provides the following ways to apply styles globally:

- StyleManager global style
- CSS global selector

e. Creating style declarations with the StyleManager

You can create CSS style declarations by using ActionScript with the `CSSStyleDeclaration` class. This lets you create and edit style sheets at run time and apply them to classes in your Flex applications.

To change the definition of the styles or to apply them during run time, you use the `setStyle()` method.

The `StyleManager` also includes a `setStyleDeclaration()` method that lets you apply a `CSSStyleDeclaration` object as a selector, so you can apply a style sheet to all components of a type. The selector can be a class or type selector.

f. Using the setStyle() and getStyle() methods

You cannot get or set style properties directly on a component as you can with other properties. Instead, you set style properties at run time by using the `getStyle()` and `setStyle()` ActionScript methods.

When you use the `getStyle()` and `setStyle()` methods, you can access the style properties of instances of objects or of style sheets. Every Flex component exposes these methods. When you are instantiating an object and setting the styles for the first time, you should try to apply style sheets rather than use the `setStyle()` method because it is computationally expensive. This method should be used only when you are changing an object's styles during run time.

g. Setting styles

The `getStyle()` method has the following signature:

```
var: return_type componentInstance.getStyle( 'property_name' )
```

The `return_type` depends on the style that you access. Styles can be of type `String`, `Number`, `Boolean`, or, in the case of skins, `Class`.

The `property_name` is a `String` that indicates the name of the style property—for example, `fontSize`.

The `setStyle()` method has the following signature:

```
componentInstance.setStyle( 'property_name', 'property_value' )
```

The `property_value` sets the new value of the specified property.

You can use the `getStyle()` method to access style properties regardless of how they were set. If you defined a style property as a tag property inline rather than in an `<mx:Style>` tag, you can get and set this style. You can override style properties that were applied in any way, such as in an `<mx:Style>` tag or in an external style sheet.

h. Loading style sheets at run time

You load a CSS-based SWF file at run time by using the `StyleManager`'s `loadStyleDeclarations()` method.

To use this method, you must import the `mx.core.StyleManager` class.

The following example shows loading a style sheet SWF file:

```
StyleManager.loadStyleDeclarations("../assets/MyStyles.swf");
```

The first parameter of the `loadStyleDeclarations()` method is the location of the style sheet SWF file to load. The location can be local or remote.

The second parameter is `update`. You set this to `true` or `false`, depending on whether you want the style sheets to immediately

update in the application.

i. Unloading style sheets at run time

You can unload a style sheet that you loaded at run time.

You do this by using the `StyleManager`'s `unloadStyleDeclarations()` method. The result of this method is that all style properties set by the specified style SWF files are returned to their defaults.

j. Using filters in Flex

You can use Adobe Flash filters to apply style-like effects to Flex components, such as Labels and Text.

You can apply filters to any visual Flex component that is derived from `UIComponent`. Filters are not styles because you cannot apply them with a style sheet or the `setStyle()` method.

The result of a filter, though, is often thought of as a style. Filters are in the `flash.filters.*` package, and include the `DropShadowFilter`, `GlowFilter`, and `BlurFilter` classes.

To apply a filter to a component with MXML, you add the filter class to the component's filters Array. The filters Array is a property inherited from the `DisplayObject` class. It contains any number of filters you want to apply to the component.

The following example applies a drop shadow to a Label control by using expanded MXML syntax and inline syntax:

<!-- Apply filter using MXML syntax to set properties.-->

```
<mx:Label text="DropShadowFilter" fontSize="20">
    <mx:filters>
        <mx:DropShadowFilter distance="10" angle="45"/>
    </mx:filters>
</mx:Label>
```

<!-- Apply filter and set properties inline. -->

```
<mx:Label text="DropShadowFilter (inline)"
    fontSize="20"
    filters="{[new DropShadowFilter(10,45)]}"/>
```

You can apply filters in ActionScript. You do this by importing the `flash.filters.*` package, and then adding the new filter to

the filters Array of the Flex control.

```
import flash.filters.*;

public function toggleFilter():void
{
    if (label1.filters.length == 0)
    {
        /*
        The first four properties of the DropShadowFilter constructor are
        distance, angle, color, and alpha.
        */
        var f:DropShadowFilter = new DropShadowFilter(5,30,0xFFFFFFFF,.8);
        var myFilters:Array = new Array();
        myFilters.push(f);

        label1.filters = myFilters;          } else {
        label1.filters = null;
    }
}
```

i) Using Blends

A value from the BlendMode class that specifies which blend mode to use.

A bitmap can be drawn internally in two ways. If you have a blend mode enabled or an external clipping mask, the bitmap is drawn by adding a bitmap-filled square shape to the vector render. If you attempt to set this property to an invalid value, Flash Player sets the value to BlendMode.NORMAL.

Flash Player applies the blendMode property on each pixel of the display object. Each pixel is composed of three constituent colors (red, green, and blue), and each constituent color has a value between 0x00 and 0xFF. Flash Player compares each constituent color of one pixel in the movie clip with the corresponding color of the pixel in the background.

For example, if blendMode is set to BlendMode.LIGHTEN, Flash Player compares the red value of the display object with the red value of the background, and uses the lighter of the two as the value for the red component of the displayed color.

The following table describes the blendMode settings. The flash.display.BlendMode class defines string values you can use. The illustrations in the table show blendMode values applied to a circular display object (2) superimposed on another display object (1).

String (BlendMode Constant Value)	Description
-------------------------------------	-------------

BlendMode.NORMAL	<p>The display object appears in front of the background. Pixel values of the display object override those of the background. Where the display object is transparent, the background is visible.</p>
BlendMode.LAYER	<p>Forces the creation of a transparency group for the display object. This means that the display object is pre-composed in a temporary buffer before it is processed further.</p> <p>This is done automatically if the display object is pre-cached using bitmap caching or if the display object is a display object container with at least one child object with a <code>blendMode</code> setting other than <code>BlendMode.NORMAL</code>.</p>
BlendMode.MULTIPLY	<p>Multiplies the values of the display object constituent colors by the colors of the background color, and then normalizes by dividing by 0xFF, resulting in darker colors. This setting is commonly used for shadows and depth effects.</p> <p>For example, if a constituent color (such as red) of one pixel in the display object and the corresponding color of the pixel in the background both have the value 0x88, the multiplied result is 0x4840.</p> <p>Dividing by 0xFF yields a value of 0x48 for that constituent color, which is a darker shade than the color of the display object or the color of the background.</p>
BlendMode.SCREEN	<p>Multiplies the complement (inverse) of the display object color by the complement of the background color, resulting in a bleaching effect. This setting is commonly used for highlights or to remove black areas of the display object.</p>

BlendMode.LIGHTEN	<p>Selects the lighter of the constituent colors of the display object and the color of the background (the colors with the larger values). This setting is commonly used for superimposing type.</p> <p>For example, if the display object has a pixel with an RGB value of 0xFFCC33, and the background pixel has an RGB value of 0xDDF800, the resulting RGB value for the displayed pixel is 0xFFFF833 (because $0xFF > 0xDD$, $0xCC < 0xF8$, and $0x33 > 0x00 = 33$).</p>
BlendMode.DARKEN	<p>Selects the darker of the constituent colors of the display object and the colors of the background (the colors with the smaller values). This setting is commonly used for superimposing type.</p> <p>For example, if the display object has a pixel with an RGB value of 0xFFCC33, and the background pixel has an RGB value of 0xDDF800, the resulting RGB value for the displayed pixel is 0xDDCC00 (because $0xFF > 0xDD$, $0xCC < 0xF8$, and $0x33 > 0x00 = 33$).</p>
BlendMode.DIFFERENCE	<p>Compares the constituent colors of the display object with the colors of its background, and subtracts the darker of the values of the two constituent colors from the lighter value. This setting is commonly used for more vibrant colors.</p> <p>For example, if the display object has a pixel with an RGB value of 0xFFCC33, and the background pixel has an RGB value of 0xDDF800, the resulting RGB value for the displayed pixel is 0x222C33 (because $0xFF - 0xDD = 0x22$, $0xF8 - 0xCC = 0x2C$, and $0x33 - 0x00 = 0x33$).</p>

BlendMode.ADD	<p>Adds the values of the constituent colors of the display object to the colors of its background, applying a ceiling of 0xFF. This setting is commonly used for animating a lightening dissolve between two objects.</p> <p>For example, if the display object has a pixel with an RGB value of 0xAAA633, and the background pixel has an RGB value of 0xDD2200, the resulting RGB value for the displayed pixel is 0xFFC833 (because $0xAA + 0xDD > 0xFF$, $0xA6 + 0x22 = 0xC8$, and $0x33 + 0x00 = 0x33$).</p>
BlendMode.SUBTRACT	<p>Subtracts the values of the constituent colors in the display object from the values of the background color, applying a floor of 0. This setting is commonly used for animating a darkening dissolve between two objects.</p> <p>For example, if the display object has a pixel with an RGB value of 0xAA2233, and the background pixel has an RGB value of 0xDDA600, the resulting RGB value for the displayed pixel is 0x338400 (because $0xDD - 0xAA = 0x33$, $0xA6 - 0x22 = 0x84$, and $0x00 - 0x33 < 0x00$).</p>
BlendMode.INVERT	Inverts the background.
BlendMode.ALPHA	<p>Applies the alpha value of each pixel of the display object to the background. This requires the <code>blendMode</code> setting of the parent display object to be set to <code>BlendMode.LAYER</code>.</p> <p>For example, in the illustration, the parent display object, which is a white background, has <code>blendMode = BlendMode.LAYER</code>.</p>
BlendMode.ERASE	<p>Erases the background based on the alpha value of the display object. This requires the <code>blendMode</code> of the parent display object to be set to <code>BlendMode.LAYER</code>. For example, in the illustration, the parent display object, which is a white background, has <code>blendMode = BlendMode.LAYER</code>.</p>

BlendMode.OVERLAY	Adjusts the color of each pixel based on the darkness of the background. If the background is lighter than 50% gray, the display object and background colors are screened, which results in a lighter color. If the background is darker than 50% gray, the colors are multiplied, which results in a darker color. This setting is commonly used for shading effects.
BlendMode.HARDLIGHT	Adjusts the color of each pixel based on the darkness of the display object. If the display object is lighter than 50% gray, the display object and background colors are screened, which results in a lighter color. If the display object is darker than 50% gray, the colors are multiplied, which results in a darker color. This setting is commonly used for shading effects.

Example code:

1.4 - Dynamically change the look of an application by using view states, transitions & effects.

a. Using View States

The properties, styles, event handlers, and child components that you define for an application or component specify its base, or default, view state.

Each view state specifies changes to the base view state, or to another view state. Consider the following when you define a view state. You can only define view states at the root of an application or at the root of a custom component; that is, as a property of the `<mx:Application>` tag of an application file, or of the root tag of an MXML component.

You define states by using the component's `states` property, normally as an `<mx:states>` tag in MXML. You populate the `states` property with an Array of one or more State objects, where each State object corresponds to a view state. You use the `name` property of the State object to specify its identifier. To change to that view state, you set a component's `currentState` property to the value of the `name` property.

a) View state overrides

A view state is defined by a set of changes, or overrides, to the base view state.

You can define the following types of overrides in a view state:

Setting the following component characteristics:

b) Properties

By using the `SetProperty` class. The following line disables the `button1` Button control as part of a view state:

```
<mx:SetProperty target="{button1}" name="enabled" value="false"/>
```

c) Styles

By using the `SetStyle` class. The following line sets the color property of the `button1` Button control as part of a view state:

```
<mx:SetStyle target="{button1}" name="color" value="0xAAAAAA"/>
```

d) Adding or removing child object

By using the `AddChild` and `RemoveChild` classes. For example, the following lines add a Button child control to the `v1` VBox control as part of a view state:

```
<mx:AddChild relativeTo="{v1}">
    <mx:Button label="New Button"/>
</mx:AddChild>
```

e) Setting or changing event handlers

By using the `SetEventHandler` class. The following line sets the click event handler of the `button1` Button control as part of a view state:

```
<mx:SetEventHandler target="{button1}" name="click"
handler="newClickHandler()"/>
```

f) Basing a view state on another view state

By default, a view state specifies a set of overrides that define the changes from the base view state to the new view state.

However, you might have a set of view states that build on each other, where the first view state defines changes relative to the base view state, and the second view state defines changes relative to the first view state, rather than to the base view state.

To define the view state relative to another view state rather than to the base view state, use the `State.basedOn` property.

When Flex changes to the new view state, it restores the base state, applies any changes from the state determined by the `basedOn` property, and then applies the changes defined in the new state.

g) Using view state events

When a component's `currentState` property changes, the `State` object for the states being exited and entered dispatch the

following events:

- `enterState`: Dispatched when a view state is entered, but not fully applied. Dispatched by a `State` object after it has been entered, and by a component after it returns to the base view state.
- `exitState`: Dispatched when a view state is about to be exited. It is dispatched by a `State` object before it is exited, and by a component before it exits the base view state.

The component on which you modify the `currentState` property to cause the state change dispatches the following events:

- `currentStateChanging`: Dispatched when the view state is about to change. It is dispatched by a component after its `currentState` property changes, but before the view state changes. You can use this event to request any data from the server required by the new view state.
- `currentStateChange`: Dispatched after the view state has completed changing. It is dispatched by a component after its `currentState` property changes. You can use this event to send data back to a server indicating the user's current view state.

h) Create a view state in `ActionScript`

In the `State` class, the `overrides` property contains an `Array` of overrides that define the view state.

The `overrides` property is the default property for the `States` class, so you can omit it in `MXML`. However, you must specify it in `ActionScript`.

Import the classes of the `mx.states` package.

Declare a `State` variable.

Create a function that does the following:

- Call the function as part of an application or custom component initialization event handler, typically in response to the `initialize` event.
- Instantiate a new `State` object.
- Assign the object to the variable you defined in Step 2.
- Set the state name.
- Declare variables for the override class objects and assign them to new instances of the classes (such as `SetProperty`).
- Specify the properties of the override instances.
- Add the override instances to the `State` object's `overrides` array.
- Add the state to the `states` array.

b. Using Transitions

a) About Transitions

View states let you vary the content and appearance of an application, typically in response to a user action. To use view states, you define the base view state of an application, and one or more additional view states.

When you change view states, Adobe Flex performs all the visual changes to the application at the same time. That means when you resize, move, or in some other way alter the appearance of the application, the application appears to jump from one view state to the next.

Instead, you might want to define a smooth visual change from one view state to the next, in which the change occurs over a period of time.

A `transition` is one or more effects grouped together to play when a view state change occurs.

For example, your application defines a form that in its base view state shows only a few fields, but in an expanded view state shows additional fields.

Rather than jumping from the base version of the form to the expanded version, you define a `Flex transition` that uses the `Resize` effect to expand the form, and then uses the `Fade` effect to slowly make the new form elements appear on the screen.

When a user changes back to the base version of the form, your transition uses a `Fade` effect to make the fields of the expanded form disappear, and then uses the `Resize` effect to slowly shrink the form back to its original size.

By using a transition, you can apply one or more effects to the form, to the fields added to the form, and to fields removed from the form. You can apply the same effects to each form field, or apply different effects.

You can also define one set of effects to play when you change state to the expanded form, and a different set of effects to play when changing from the expanded state back to the base state.

b) Defining Transitions

You use the `Transition` class to create a transition.

The following table defines the properties of the `Transition` class:

Property	Definition
<code>fromState</code>	<p>A <code>String</code> that specifies the view state that you are changing from when you apply the transition.</p> <p>The default value is an asterisk, "*", which means any view state.</p>
<code>toState</code>	<p>A <code>String</code> that specifies the view state that you are changing to when you apply the transition.</p> <p>The default value is an asterisk, "*", which means any view state.</p>

effect	<p>The Effect object to play when you apply the transition.</p> <p>Typically, this is a composite effect, such as the <code>Parallel</code> or <code>Sequence</code> effect that contains multiple effects.</p>
--------	---

You can define multiple `Transition` objects in an application. The `UIComponent` class includes a `transitions` property that you set to an Array of `Transition` objects.

You define the transitions for this example in MXML using the `<mx:transitions>` tag, as the following example shows:

```
<!-- Define Transition array with one Transition object.-->
<mx:transitions>
  <!-- A transition for changing from any state to any state. -->
  <mx:Transition id="myTransition" fromState="*" toState="*">
    <!-- Define a Parallel effect as the top-level effect.-->
    <mx:Parallel id="t1" targets="{ [p1,p2,p3] }">
      <!-- Define a Move and Resize effect.-->
      <mx:Move duration="400"/>
      <mx:Resize duration="400"/>
    </mx:Parallel>
  </mx:Transition>
</mx:transitions>
```

c) Defining Multiple Transitions

You can define multiple transitions in your application so that you can associate a specific transition with a specific change to the view state. To specify the transition associated with a change to the view states, you use the `fromState` and `toState` properties.

By default, both the `fromState` and `toState` properties are set to `"*"`, which indicates that the transition should be applied to any changes in the view state. You can set either property to an empty string, `""`, which corresponds to the base view state.

You use the `fromState` property to explicitly specify the view state that your are changing from, and the `toState` property to explicitly specify the view state that you are changing to, as the following example shows:

```
<mx:transitions>
  <!-- Play for a change to the login view state from any view state. -->
  <mx:Transition id="toLoginFromAny" fromState="*" toState="login">
    ...
  </mx:Transition>

  <!-- Play for a change to the login view state from the details view state. -->
  <mx:Transition id="toLoginFromDetails" fromState="details" toState="login">
    ...
  </mx:Transition>
```

```

    <!-- Play for a change from any view state to any other view state. -->
    <mx:Transition id="toAnyFromAny" fromState="*" toState="*">
        ...
    </mx:Transition>
</mx:transitions>

<!-- Go to the login view state, transition toLoginFromAny plays. -->
<mx:Button click="currentState='login';"/>

<!-- Go to the details view state, transition toAnyFromAny plays. -->
<mx:Button click="currentState='details';"/>

<!-- Go to the login view state, transition toLoginFromDetails plays because you
transitioned from the details to the login view state. -->
<mx:Button click="currentState='login';"/>

<!-- Go to the base view state, transition toAnyFromAny plays. -->
<mx:Button click="currentState='';"/>

```

If a state change matches two transitions, the `toState` property takes precedence over the `fromState` property. If more than one transition matches, Flex uses the first matching definition detected in the transition Array.

d) Defining Effect Targets

The `<mx:Transition>` tag shown in the section defining transitions defines the effects that make up a transition. The top-level effect defines the target components of the effects in the transition when the effect does not explicitly define a target.

In that example, the transition is performed on all three Panel containers in the application. If you want the transition to play only on the first two panels, you define the Parallel effects as the following example shows:

```

<mx:Transition fromState="*" toState="*">
    <!-- Define a Parallel effect as the top most effect.-->
    <mx:Parallel id="t1" targets="{ [p1,p2] }">
        <mx:Move duration="400"/>
        <mx:Resize duration="400"/>
    </mx:Parallel>
</mx:Transition>

```

You removed the third panel from the transition, so it is no longer a target of the Move and Resize effects. Therefore, the third panel appears to jump to its new position and size during the change in view state.

The other two panels show a smooth change in size and position for the 400-millisecond (ms) duration of the effects.

You can also use the `target` or `targets` properties of the effects within the transition to explicitly specify the effect target, as the following example shows:

```

<mx:Parallel id="t1" targets="{ [p1,p2,p3] }">
    <mx:Move targets="{ [p1,p2] }" duration="400"/>
    <mx:Resize duration="400"/>
</mx:Parallel>

```

In this example, the `Resize` effect plays on all three panels, while the `Move` effect plays only on the first two panels. You could also write this example as the following code shows:

```
<mx:Parallel id="t1" targets="{[p1,p2]}">
    <mx:Move duration="400"/>
    <mx:Resize targets="{ [p1,p2,p3] }" duration="400"/>
</mx:Parallel>
```

e) Defining the effect start and end values

Like any effect, an effect within a transition has properties that you use to configure it.

For example, most effects have properties that define starting and ending information for the target component, such as the `xFrom`, `yFrom`, `xTo`, and `yTo` properties of the `Move` effect.

Effects defined in a transition must determine their property values for the effect to execute. Flex uses the following rules to determine the start and end values of effect properties of a transition:

1. If the effect explicitly defines the values of any properties, use them in the transition, as the following example shows:

```
<mx:Transition fromState="*" toState="*">
<mx:Sequence id="t1" targets="{ [p1,p2,p3] }">
    <mx:Blur duration="100"
        blurXFrom="0.0"
        blurXTo="10.0"
        blurYFrom="0.0"
        blurYTo="10.0"/>
    <mx:Parallel>
        <mx:Move duration="400"/>
        <mx:Resize duration="400"/>
    </mx:Parallel>
</mx:Sequence>
</mx:Transition>
```

In this example, the two `Blur` filters explicitly define the properties of the effect.

If the effect does not explicitly define the start values of the effect, Flex determines them from the current settings of the target component, as defined by the current view state.

In the example in rule 1, notices that the `Move` and `Resize` effects do not define start values. Therefore, Flex determines them from the current size and position of the effect targets in the current view state.

If the effect does not explicitly define the end values of the effect, Flex determines them from the settings of the target component in the destination view state.

In the example in rule 1, the `Move` and `Resize` effects determine the end values from the size and position of the effect targets in the destination view state. In some cases, the destination view state explicitly defines these values. If the destination view state does not define the values, Flex determines them from the settings of the base view state.

If there are no explicit values, and Flex cannot determine values from the current or destination view states, the effect uses its default property values.

f) Handling events when using transitions

You can handle view state events, such as `currentStateChange` and `currentStateChanging`, as part of an application that defines transitions.

Changes to the view state, dispatching events, and playing transition effects occur in the following order:

1. You set the `currentState` property to the destination view state.
2. Flex dispatches the `currentStateChanging` event.
3. Flex examines the list of transitions to determine the one that matches the change of the view state.
4. Flex examines the components to determine the start values of any effects.
5. Flex applies the destination view state to the application.
6. Flex dispatches the `currentStateChange` event.
7. Flex plays the effects that you defined in the transition.

If you change state again while a transition plays, Flex jumps to the end of the transition before starting any transition associated with the new change of view state.

g) Using action effects in a transition

To move from the base view state to the `OneOnly` view state, you create the following view state definition:

```
<mx:states>
  <mx:State name="OneOnly">
    <mx:SetProperty target="{p2}" name="visible" value="false"/>
    <mx:SetProperty target="{p2}" name="includeInLayout" value="false"/>
  </mx:State>
</mx:states>
```

You set the value of the `visible` and `includeInLayout` properties to `false` so that Flex makes the second Panel container invisible and ignores it when laying out the application.

If the `visible` property is `false`, and the `includeInLayout` property is `true`, the container is invisible, but Flex lays out the application as if the component were visible.

A view state defines how to change states, and the transition defines the order in which the visual changes occur.

In the you play an Iris effect on the second panel when it disappears, and when it reappears on a transition back to the base state.

For the change from the base state to the `OneOnly` state, you define the `toOneOnly` transition which uses the Iris effect to make the second panel disappear, and then sets the panel's `visible` and `includeInLayout` properties to `false`.

For a transition back to the base state, you define the `toAnyFromAny` transition that makes the second panel visible by setting its `visible` and `includeInLayout` properties to `true`, and then uses the Iris effect to make the panel appear, as the following example shows:

```
<mx:transitions>
  <mx:Transition id="toOneOnly" fromState="*" toState="OneOnly">
    <mx:Sequence id="t1" targets="{ [p2] }">
      <mx:Iris showTarget="false" duration="350"/>
      <mx:SetPropertyAction name="visible"/>
      <mx:SetPropertyAction target="{p2}" name="includeInLayout"/>
    </mx:Sequence>
  </mx:Transition>

  <mx:Transition id="toAnyFromAny" fromState="*" toState="*">
    <mx:Sequence id="t2" targets="{ [p2] }">
      <mx:SetPropertyAction target="{p2}" name="includeInLayout"/>
      <mx:SetPropertyAction name="visible"/>
      <mx:Iris showTarget="true" duration="350"/>
    </mx:Sequence>
  </mx:Transition>
</mx:transitions>
```

In the `toOneOnly` transition, if you hide the target by setting its `visible` property to `false`, and then play the Iris effect, you would not see the Iris effect play because the target is already invisible.

To control the order of view state changes during a transition, Flex defines several action effects. The previous example uses the `SetPropertyAction` action effect to control when to set the `visible` and `includeInLayout` properties in the transition.

The following table describes the action effects:

Action Effect	Corresponding view state class	Use
<code>SetPropertyAction</code>	<code>SetProperty</code>	Sets a property value as part of a transition.

SetStyleAction	SetStyle	Sets a style to a value as part of a transition.
AddChildAction	AddChild	Adds a child as part of a transition.
RemoveChildAction	RemoveChild	Removes a child as part of a transition.

When you create a view state, you use the `SetProperty`, `SetStyle`, `AddChild`, and `RemoveChild` classes to define the view state.

To control when a change defined by the view state property occurs in a transition, you use the corresponding action effect. The action effects give you control over the order of the state change.

In the previous example, you used the following statement to define an action effect to occur when the value of the visible property of a component changes:

```
<mx:SetPropertyAction name="visible"/>
```

This action effect plays when the value of the visible property changes to either true or false.

You can further control the effect using the value property of the `<mx:SetPropertyAction>` tag, as the following example shows:

```
<mx:SetPropertyAction name="visible" value="true"/>
```

In this example, you specify to play the effect only when the value of the visible property changes to true. Adding this type of information to the action effect can be useful if you want to use filters with your transitions. The action effects do not support a duration property; they only perform the specified action.

h) Filtering effects

By default, Flex applies all of the effects defined in a transition to all of the target components of the transition.

Therefore, in the following example, Flex applies the `Move` and `Resize` effects to all three targets:

```
<mx:Transition fromState="*" toState="*">

    <!-- Define a Parallel effect as the top most effect.-->
    <mx:Parallel id="t1" targets="{ [p1,p2,p3] }">
        <!-- Define a Move and Resize effect.-->
        <mx:Move duration="400"/>
        <mx:Resize duration="400"/>
    </mx:Parallel>
</mx:Transition>
```

However, you might want to conditionalize an effect so that it does not apply to all target components, but only to a subset of the components.

Each change of view state removes the top panel, moves the bottom panel to the top, and adds the next panel to the bottom of the screen. In this example, the third panel is invisible in the base view state.

For this example, you define a single transition that applies a `WipeUp` effect to the top panel as it is removed, applies a `Move` effect to the bottom panel as it moves to the top position, and applies another `WipeUp` effect to the panel being added to the bottom, as the following example shows:

```
<mx:transitions>
  <mx:Transition fromState="*" toState="*">
    <mx:Sequence targets="{ [p1,p2,p3] }">
<mx:Sequence id="sequence1" filter="hide">
<mx:WipeUp/>

      <mx:SetPropertyAction name="visible" value="false"/>
    </mx:Sequence>
    <mx:Move filter="move"/>
    <mx:Sequence id="sequence2" filter="show">
      <mx:SetPropertyAction name="visible" value="true"/>
      <mx:WipeUp/>
    </mx:Sequence>
  </mx:Sequence>
</mx:Transition>
</mx:transitions>
```

The `sequence1` `Sequence` effect uses the `filter` property to specify the change that a component must go through in order for the effect to play on it. In this example, the `sequence1` effect specifies a value of `"hide"` for the `filter` property.

Therefore, the `WipeUp` and `SetPropertyAction` effects only play on those components that change from visible to invisible by setting their `visible` property to `false`.

In the `sequence2` `Sequence` effect, you set the `filter` property to `show`. Therefore, the `WipeUp` and `SetPropertyAction` effects only play on those components whose state changes from invisible to visible by setting their `visible` property to `true`.

The `Move` effect also specifies a `filter` property with a value of `move`. Therefore, it applies to all target components that are changing position.

The following table describes the possible values of the `filter` property:

Value	Description
-------	-------------

Add	Specifies to play the effect on all children added during the change of view state.
Hide	Specifies to play the effect on all children whose visible property changes from true to false during the change of view state.
Move	Specifies to play the effect on all children whose x or y properties change during the change of view state.
Remove	Specifies to play the effect on all children removed during the change of view state.
Resize	Specifies to play the effect on all children whose width or height properties change during the change of view state.
Show	Specifies to play the effect on all children whose visible property changes from false to true during the change of view state.

sing Effects

1.5 - Position UI elements by using constraint-based layout.

a) Absolute positioning

Three containers support `absolute` positioning:


Application and Panel controls use absolute positioning if you specify the layout property as "`absolute`" (`ContainerLayout.ABSOLUTE`).

The Canvas container always uses `absolute` positioning.

With absolute positioning, you specify the child control position by using its `x` and `y` properties, or you specify a constraint-based layout; otherwise, Flex places the child at position 0, 0 of the parent container.

When you specify the `x` and `y` coordinates, Flex repositions the controls only when you change the property values.

When you use absolute positioning, you have full control over the locations of the container's children. This lets you overlap components.

 Position UI elements by using enhanced constraints.

1.7 - Implement application navigation by using navigator containers.

a. ViewStack

- `selectedIndex`: The index of the currently active container if one or more child containers are defined. The value of this property is `-1` if no child containers are defined.

The index is numbered from `0` to `numChildren - 1`, where `numChildren` is the number of child containers in the `ViewStack` container. Set this property to the index of the container that you want active.

You can use the `selectedIndex` property of the `<mx:ViewStack>` tag to set the default active container when your application starts.

- `selectedChild`: The currently active container if one or more child containers are defined. The value of this property is `null` if no child containers are defined.

Set this property in ActionScript to the identifier of the container that you want active

You can set this property only in an ActionScript statement, not in MXML.

The following example uses ActionScript to set the `selectedChild` property so that the active child container is the child container with an identifier of `search`:

- `numChildren`: Contains the number of child containers in the `ViewStack` container.

The following uses the `numChildren` property in an ActionScript statement to set the active child container to the last container in the stack:

```
myViewStack.selectedIndex = myViewStack.numChildren - 1;
```

b. Tab Navigator

A `TabNavigator` container creates and manages a set of tabs, which you use to navigate among its children. The children of a `TabNavigator` container are other containers.

The `TabNavigator` container creates one tab for each child. When the user selects a tab, the `TabNavigator` container displays the associated child,

You can also set the currently active child by using the `selectedChild` and `selectedIndex` properties inherited from the `ViewStack` container as follows:

- `numChildren`: Contains the number of child containers in the `ViewStack` container.

The following uses the `numChildren` property in an `ActionScript` statement to set the active child container to the last container in the stack:

```
myViewStack.selectedIndex = myViewStack.numChildren - 1;
```

- `selectedChild`: The currently active container if one or more child containers are defined. The value of this property is null if no child containers are defined.

Set this property in `ActionScript` to the identifier of the container that you want active

You can set this property only in an `ActionScript` statement, not in `MXML`.

The following example uses `ActionScript` to set the `selectedChild` property so that the active child container is the child container with an identifier of `search`:

For more information on the `selectedChild` and `selectedIndex` properties, see `ViewStack` navigator container.

c. Accordion

An `Accordion` navigator container has a collection of child containers, but only one of them at a time is visible. It creates and manages navigator buttons (accordion headers), which you use to navigate between the children.

There is one navigator button associated with each child container, and each navigator button belongs to the `Accordion` container, not to the child. When the user clicks a navigator button, the associated child container is displayed.

The transition to the new child uses an animation to make it clear to the user that one child is disappearing and a different one is appearing.

The `Accordion` container does not extend the `ViewStack` container, but it implements all the properties, methods, styles, and events of the `ViewStack` container, such as `selectedIndex` and `selectedChild`.

1.8 - Customize list-based controls.

a. List Control

The `List` control displays a vertical list of items. Its functionality is very similar to that of the `SELECT` form element in HTML.

If there are more items than can be displayed at once, it can display a vertical scroll bar so the user can access all items in the list.

An optional horizontal scroll bar lets the user view items when the full width of the list items is unlikely to fit. The user can select one or more items from the list, depending on the value of the `allowMultipleSelection` property.

Note: The `HorizontalList`, `TileList`, `DataGrid`, `Menu`, and `Tree` controls are derived from the `List` control or its immediate parent, the `ListBase` class. As a result, much of the information for the `List` control applies to these controls.

b. Horizontal List Control

The `HorizontalList` control displays a horizontal list of items.

If there are more items than can be displayed at once, it can display a horizontal scroll bar so the user can access all items in the list.

c. Tile List Control

The `TileList` control displays a number of items laid out in tiles.

It displays a scroll bar on one of its axes to access all items in the list, depending on the `direction` property. You can set the size of the tiles by using `rowHeight` and `columnWidth` properties.

Alternatively, Flex measures the item renderer for the first item in the `dataProvider` and uses that size for all tiles.

d. ComboBox Control

e. DataGrid Control

The `DataGrid` control is like a `List` except that it can show more than one column of data making it suited for showing objects with multiple properties.

The `DataGrid` control provides the following features:

- Columns of different widths or identical fixed widths
- Columns that the user can resize at runtime

- Columns that the user can reorder at runtime
- Optional customizable column headers
- Ability to use a custom item renderer for any column to display data other than text
- Support for sorting the data by clicking on a column

The `DataGrid` control is intended for viewing data, and not as a layout tool like an `HTML table`. The `mx.containers` package provides those layout tools.

f. Tree Control

The `Tree` control lets a user view hierarchical data arranged as an expandable tree.

Each item in a tree can be a leaf or a branch. A leaf item is an end point in the tree. A branch item can contain leaf or branch items, or it can be empty.

By default, a leaf is represented by a text label next to a file icon.

A branch is represented by a text label next to a folder icon, with a disclosure triangle that a user can open to expose children.

The `Tree` class uses an `ITreeDataDescriptor` or `ITreeDataDescriptor2` object to parse and manipulate the data provider.

The default tree data descriptor, an object of the `DefaultDataDescriptor` class, supports `XML` and `Object` classes; an `Object` class data provider must have all children in `children` fields.

The `Tree` control has the following default sizing characteristics:

Characteristic	Description
Default Size	Wide enough to accommodate the icon, label, and expansion triangle, if any, of the widest node in the first 7 displayed (uncollapsed) rows, and seven rows high, where each row is 20 pixels in height. If a scroll bar is required, the width of the scroll bar is not included in the width calculations.
Minimum size	0 pixels.
Maximum size	5000 by 5000

2. Flex system architecture and design

2.1 Create and use custom components.

Your application might require you to create components, rather than modifying existing ones.

To create components, you typically create them in `ActionScript` by creating a subclass from the `UIComponent` class. This class contains the generic functionality of all Flex components.

You then add the required functionality to your new component to meet your application requirements.

a. User-interface, or visual, components

User-interface components contain both processing logic and visual elements.

You create custom user-interface components to modify existing behavior or add new functionality to the component. These components usually extend the Flex component hierarchy.

You can extend from the `UIComponent` class, or any of the Flex components, such as `Button`, `ComboBox`, or `DataGrid`.

Your custom `ActionScript` component inherits all of the methods, properties, events, styles, and effects of its superclass.

b. Non-visual components

Non-visual components define non-visual elements.

Flex includes several types of non-visual components that you can create, including `formatters`, `validators`, and `effects`.

You create non-visual components by creating a subclass from the Flex component hierarchy.

- For `validators`, you create subclasses of the `Validator` class
- For `formatters` you create subclasses of the `Formatter` class
- For `effects`, you create subclasses of the `Effect` class

2.3 Transfer data between components by using data bindings.

a. Using Databinding

Data binding is the process of tying the data in one object to another object. It provides a convenient way to pass data between the different layers of the application. Data binding requires a source property, a destination property, and a triggering event that indicates when to copy the data from the source to the destination.

An object dispatches the triggering event when the source property changes. Adobe Flex provides three ways to specify data

binding: the curly braces (`{ }`) syntax in MXML, the `<mx:Binding>` tag in MXML, and the `BindingUtils` methods in `ActionScript`.

a) Defining data bindings in `ActionScript`

You cannot include `ActionScript` code in a data binding expression defined by the `bindProperty()` or `bindSetter()` method. Instead, use the `bindSetter()` method to specify a method to call when the binding occurs.

You cannot include an `E4X` expression in a data binding expression defined in `ActionScript`.

You cannot include functions or array elements in property chains in a data binding expression defined by the `bindProperty()` or `bindSetter()` method.

The MXML compiler has better warning and error detection support than runtime data bindings defined by the `bindProperty()` or `bindSetter()` method.

You can define a data binding in `ActionScript` by using the `mx.binding.utils.BindingUtils` class. This class defines static methods that let you create a data binding to a property implemented as a variable, by using the `bindProperty()` method, or to a method, by using the `bindSetter()` method.

b) Differences between defining bindings in MXML and `ActionScript`

There are a few differences between defining data bindings in MXML at compile time and in defining them at runtime in `ActionScript`:

The arguments to the `bindSetter()` method specify the following:

The source object

The name of the source property

A method that is called when the source property changes

2.4 Create, Handle, and Dispatch Custom Events.

a. Using an event object

When a Flex component dispatches an event, it creates an event object, where the properties of the event object contain information describing the event. An event listener takes this event object as an argument and accesses the properties of the object to determine information about the event.

The base class for all event objects is the `flash.events.Event` class. All event objects are instances of the `Event` class, or instances of a subclass of the `Event` class.

The following table describes the public properties of the Event class. The Event class implements these properties using getter methods.

Property	Type	Description
type	String	The name of the event; for example, "click". The event constructor sets this property.
target	EventDispatcher	A reference to the component instance that dispatches the event. This property is set by the <code>dispatchEvent()</code> method; you cannot change this to a different object.
currentTarget	EventDispatcher	A reference to the component instance that is actively processing the Event object. The value of this property is different from the value of the target property during the event capture and bubbling phase.
eventPhase	uint	The current phase in the event flow. The property might contain the following values: <code>EventPhase.CAPTURING_PHASE</code> : The capture phase <code>EventPhase.AT_TARGET</code> : The target phase <code>EventPhase.BUBBLING_PHASE</code> : The bubbling phase
bubbles	Boolean	Whether an event is a bubbling event. If the event can bubble, the value for this property is true; otherwise, it is false. You can optionally pass this property as a constructor argument to the Event class. By default, most event classes set this property to false.
cancelable	Boolean	Whether the event can be canceled. If the event can be canceled, the value for this value is true; otherwise, it is false. You can optionally pass this property as a constructor argument to the Event class. By default, most event classes set this property to false.

2.5 Dispatching Custom Events

Create a subclass from the `flash.events.Event` class to create an event class that describes the event object.

Flex defines many of the most common events, such as the click event for the Button control; however, your application may require that you create events. In your custom Flex components, you can dispatch any of the predefined events inherited by the component from its superclass, and dispatch new events that you define within the component.

To dispatch a new event from your custom component, you must do the following:

(Optional) Use the `[Event]` metadata tag to make the event public so that the MXML compiler recognizes it.

Dispatch the event using the `dispatchEvent()` method.

You use the `dispatchEvent()` method to dispatch an event. The `dispatchEvent()` method has the following signature:

```
public dispatchEvent(event:Event):Boolean dispatchEvent(new Event("click"));
```

This method requires an argument of the `Event` type, which is the event object. The `dispatchEvent()` method initializes the `target` property of the event object with a reference to the component dispatching the event.

You can create an event object and dispatch the event in a single statement, as the following example shows:

```
dispatchEvent(new Event("click"));
```

You can also create an event object, initialize it, and then dispatch it, as the following example shows:

```
var eventObj:EnableChangeEvent = new EnableChangeEvent("enableChange");
    eventObj.isEnabled = true;
dispatchEvent(eventObj);
```

a. Creating static constants for the `Event.type` property

The constructor of an event class typically takes a single required argument that specifies the value of the event object's `type` property. In the previous section, you passed the string `enableChange` to the constructor, as the following example shows:

```
// Define event object, initialize it, then dispatch it.
var eventObj:EnableChangeEvent = new EnableChangeEvent("enableChange");
    dispatchEvent(eventObj);
```

2.6 Handle framework events.

2.7 List and describe the differences between model, view, and controller code in a Flex application.

MVC is a design pattern or software architecture that separates the applications data, user interface, and control logic into three distinct groupings. The goal is to implement the logic so changes can be made to one portion of the application with a minimal impact to the others. Short definitions of the key terms are as follows:

- **Model:** The data the application uses. It manages data elements, responds to queries about its state, and instructions to change the data.
- **View:** The user interface. It is responsible for presenting model data to the user and gathering information from the user.
- **Controller:** Responds to events, typically user events, but also system events. The events are interpreted and the controller invokes changes on the model and view.

Generally, the flow of **MVC** is as follows:

- User interacts with the user interface (a view), such as clicking a button to add an item to a shopping cart.
- The controller handles the input event
- The controller accesses the model, maybe by retrieving or altering data.
- The view then uses the model data for appropriate user presentation.

3. Programming Flex applications with ActionScript

3.1 Define and extend an ActionScript class.

a. Classes

A class is an abstract representation of an object. A class stores information about the types of data that an object can hold and the behaviors that an object can exhibit. The usefulness of such an abstraction may not be apparent when you write small scripts that contain only a few objects interacting with one another. As the scope of a program grows, however, and the number of objects that must be managed increases, you may find that classes allow you to better control how objects are created and how they interact with one another.

As far back as ActionScript 1.0, ActionScript programmers could use Function objects to create constructs that resembled classes. ActionScript 2.0 added formal support for classes with keywords such as `class` and `extends`. ActionScript 3.0 not only continues to support the keywords introduced in ActionScript 2.0, but also adds some new capabilities, such as enhanced access control with the `protected` and `internal` attributes, and better control over inheritance with the `final` and `override` keywords.

b. Class definitions

ActionScript 3.0 class definitions use syntax that is similar to that used in ActionScript 2.0 class definitions. Proper syntax for a class definition calls for the `class` keyword followed by the class name.

The class body, which is enclosed by curly braces (`{ }`), follows the class name.

For example, the following code creates a class named `Shape` that contains one variable, named `visible`:

```
public class Shape
{
    var visible:Boolean = true;
}
```

c. Class attributes

ActionScript 3.0 allows you to modify class definitions using one of the following four attributes:

Attribute	Definition
dynamic	Allow properties to be added to instances at run time.
final	Must not be extended by another class.
internal (default)	Visible to references inside the current package
public	Visible to references everywhere

For each of these attributes, except for internal, you must explicitly include the attribute to get the associated behavior.

For example, if you do not include the dynamic attribute when defining a class, you will not be able to add properties to a class instance at run time.

You explicitly assign an attribute by placing it at the beginning of the class definition, as the following code demonstrates:

```
dynamic class Shape
{
    //...
}
```

Notice that the list does not include an attribute named `abstract`. This is because abstract classes are not supported in ActionScript 3.0.

Notice also that the list does not include attributes named `private` and `protected`. These attributes have meaning only inside a class definition, and cannot be applied to classes themselves. If you do not want a class to be publicly visible outside a package, place the class inside a package and mark the class with the internal attribute.

Alternatively, you can omit both the internal and public attributes, and the compiler will automatically add the internal attribute for you. If you do not want a class to be visible outside the source file in which it is defined, place the class at the bottom of your source file, below the closing curly brace of the package definition.

d. Class body

The class body, which is enclosed by curly braces, is used to define the variables, constants, and methods of your class. The following example shows the declaration for the Accessibility class in the Adobe Flash Player API:

```
public final class Accessibility
{
    public static function get active():Boolean;
    public static function updateProperties():void;
}
```

You can also define a namespace inside a class body. The following example shows how a namespace can be defined within a class body and used as an attribute of a method in that class:

```
public class SampleClass
{
    public namespace sampleNamespace;
    sampleNamespace function doSomething():void;
}
```

ActionScript 3.0 allows you to include not only definitions in a class body, but also statements. Statements that are inside a class body, but outside a method definition, are executed exactly once—when the class definition is first encountered and the associated class object is created.

The following example includes a call to an external function, `hello()`, and a trace statement that outputs a confirmation message when the class is defined:

```
function hello():String
{
    trace("hola");
}

class SampleClass
{
    hello();
    trace("class created");
}
// output when class is created
hola
class created
```

In contrast to previous versions of ActionScript, in ActionScript 3.0 it is permissible to define a static property and an instance property with the same name in the same class body.

For example, the following code declares a static variable named `message` and an instance variable of the same name:

```
class StaticTest
{
    static var message:String = "static variable";
    var message:String = "instance variable";
}

// In your script
var myST:StaticTest = new StaticTest();
trace(StaticTest.message); // output: static variable
```

```
trace(myST.message); // output: instance variable
```

e. Class property attributes

In discussions of the ActionScript object model, the term property means anything that can be a member of a class, including variables, constants, and methods. This differs from the way the term is used in the ActionScript 3.0

Language and Components Reference, where the term is used more narrowly and includes only class members that are variables or are defined by a getter or setter method. In ActionScript 3.0, there is a set of attributes that can be used with any property of a class. The following table lists this set of attributes.

Attribute	Definition
<code>internal</code> (default)	Visible to references inside the current package
<code>private</code>	Visible to references in the same class
<code>protected</code>	Visible to references in the same class and derived classes
<code>public</code>	Visible to references everywhere
<code>static</code>	Specifies that a property belongs to the class, as opposed to instances of the class.
<code>UserDefinedNamespace</code>	Custom namespace name defined by user

f. Access control namespace attributes

ActionScript 3.0 provides four special attributes that control access to properties defined inside a class: `public`, `private`, `protected`, and `internal`.

The `public` attribute makes a property visible anywhere in your script.

For example, to make a method available to code outside its package, you must declare the method with the `public` attribute. This is true for any property, whether it is declared using the `var`, `const`, or `function` keywords.

The `private` attribute makes a property visible only to callers within the property's defining class. This behavior differs from that of the `private` attribute in ActionScript 2.0, which allowed a subclass to access a private property in a superclass. Another significant change in behavior has to do with run-time access.

In ActionScript 2.0, the `private` keyword prohibited access only at compile time and was easily circumvented at run time.

In ActionScript 3.0, this is no longer true. Properties that are marked as `private` are unavailable at both compile time and run time.

For example, the following code creates a simple class named PrivateExample with one private variable, and then attempts to access the private variable from outside the class.

In ActionScript 2.0, compile-time access was prohibited, but the prohibition was easily circumvented by using the property access operator ([]), which does the property lookup at run time rather than at compile time.

```
class PrivateExample
{
    private var privVar:String = "private variable";
}

var myExample:PrivateExample = new PrivateExample();

trace(myExample.privVar); // compile-time error in strict mode

trace(myExample["privVar"]); // ActionScript 2.0 allows access,
//but in ActionScript 3.0, this is a run-time error.
```

In ActionScript 3.0, an attempt to access a private property using the dot operator (myExample.privVar) results in a compile-time error if you are using strict mode. Otherwise, the error is reported at run time, just as it is when you use the property access operator (myExample["privVar"]).

The following table summarizes the results of attempting to access a private property that belongs to a sealed (not dynamic) class:

	Strict Mode	Standard Mode
Dot operator (.)	Compile-time error	Run-time error
Bracket operator ([])	Run-time error	Run-time error

In classes declared with the dynamic attribute, attempts to access a private variable will not result in a run-time error. Instead, the variable is simply not visible, so Flash Player or Adobe AIR returns the value undefined.

A compiletime error occurs, however, if you use the dot operator in strict mode.

The following example is the same as the previous example, except that the PrivateExample class is declared as a dynamic class:

```
dynamic class PrivateExample
{
    private var privVar:String = "private variable";
}

var myExample:PrivateExample = new PrivateExample();
trace(myExample.privVar); // compile-time error in strict mode
trace(myExample["privVar"]); // output: undefined
```

g. static attribute

The static attribute, which can be used with properties declared with the `var`, `const`, or function keywords, allows you to attach a property to the class rather than to instances of the class.

Code external to the class must call static properties by using the class name instead of an instance name. Static properties are not inherited by subclasses, but the properties are part of a subclass's scope chain.

This means that within the body of a subclass, a static variable or method can be used without referencing the class in which it was defined.

h. User-defined namespace attributes

As an alternative to the predefined access control attributes, you can create a custom namespace for use as an attribute.

Only one namespace attribute can be used per definition, and you cannot use a namespace attribute in combination with any of the access control attributes (`public`, `private`, `protected`, `internal`).

i. Variables

Variables can be declared with either the `var` or `const` keywords. Variables declared with the `var` keyword can have their values changed multiple times throughout the execution of a script.

Variables declared with the `const` keyword are called `constants`, and can have values assigned to them only once. An attempt to assign a new value to an initialized constant results in an error.

j. Static variables

Static variables are declared using a combination of the `static` keyword and either the `var` or `const` statement.

Static variables, which are attached to a class rather than an instance of a class, are useful for storing and sharing information that applies to an entire class of objects.

For example, a static variable is appropriate if you want to keep a tally of the number of times a class is instantiated or if you want to store the maximum number of class instances that are allowed.

The following example creates a `totalCount` variable to track the number of class instantiations and a `MAX_NUM` constant to store the maximum number of instantiations.

The `totalCount` and `MAX_NUM` variables are static, because they contain values that apply to the class as a whole rather than to a particular instance.

```
class StaticVars
{
    public static var totalCount:int = 0;
    public static const MAX_NUM:uint = 16;
}
```

Code that is external to the `StaticVars` class and any of its subclasses can reference the `totalCount` and `MAX_NUM`

properties only through the class itself.

For example, the following code works:

```
trace(StaticVars.totalCount); // output: 0
trace(StaticVars.MAX_NUM); // output: 16
```

You cannot access static variables through an instance of the class, so the following code returns errors:

```
var myStaticVars:StaticVars = new StaticVars();
trace(myStaticVars.totalCount); // error
trace(myStaticVars.MAX_NUM); // error
```

Variables that are declared with both the `static` and `const` keywords must be initialized at the same time as you declare the constant, as the `StaticVars` class does for `MAX_NUM`.

You cannot assign a value to `MAX_NUM` inside the constructor or an instance method. The following code will generate an error, because it is not a valid way to initialize a static constant:

k. Instance variables

Instance variables include properties declared with the `var` and `const` keywords, but without the `static` keyword. Instance variables, which are attached to class instances rather than to an entire class, are useful for storing values that are specific to an instance.

For example, the `Array` class has an instance property named `length`, which stores the number of array elements that a particular instance of the `Array` class holds.

Instance variables, whether declared as `var` or `const`, cannot be overridden in a subclass. You can, however, achieve functionality that is similar to overriding variables by overriding `getter` and `setter` methods.

3.2 Implement an ActionScript interface.

An interface is a collection of method declarations that allows unrelated objects to communicate with one another.

For example, the Flash Player API defines the `IEventDispatcher` interface, which contains method declarations that a class can use to handle event objects.

The `IEventDispatcher` interface establishes a standard way for objects to pass event objects to one another.

The following code shows the definition of the `IEventDispatcher` interface:

```
public interface IEventDispatcher
{
    function addEventListener( type:String,
                             listener:Function,
                             useCapture:Boolean=false,
                             priority:int=0,
                             useWeakReference:Boolean = false):void;
```



```

function removeEventListener(type:String,
                             listener:Function,
                             useCapture:Boolean=false) :void;

function dispatchEvent(event:Event) :Boolean;
function hasEventListener(type:String) :Boolean;
function willTrigger(type:String) :Boolean;
}

```

Interfaces are based on the distinction between a method's interface and its implementation. A method's interface includes all the information necessary to invoke that method, including the name of the method, all of its parameters, and its return type.

A method's implementation includes not only the interface information, but also the executable statements that carry out the method's behavior.

An interface definition contains only method interfaces, and any class that implements the interface is responsible for defining the method implementations.

a. Defining an interface

The structure of an interface definition is similar to that of a class definition, except that an interface can contain only methods with no method bodies.

Interfaces cannot include variables or constants but can include `getters` and `setters`.

To define an interface, use the `interface` keyword.

For example, the following interface, `IExternalizable`, is part of the `flash.utils` package in the Flash Player API.

The `IExternalizable` interface defines a protocol for serializing an object, which means converting an object into a format suitable for storage on a device or for transport across a network.

```

public interface IExternalizable
{
    function writeExternal(output:IDataOutput) :void;
    function readExternal(input:IDataInput) :void;
}

```

The Flash Player API follows a convention in which interface names begin with an uppercase `I`, but you can use any legal identifier as an interface name. Interface definitions are often placed at the top level of a package.

Interface definitions cannot be placed inside a class definition or inside another interface definition. Interfaces can extend one or

more other interfaces.

For example, the following interface, `IExample`, extends the `IExternalizable` interface:

```
public interface IExample extends IExternalizable
{
    function extra():void;
}
```

Any class that implements the `IExample` interface must include implementations not only for the `extra()` method, but also for the `writeExternal()` and `readExternal()` methods inherited from the `IExternalizable` interface.

b. Implementing an interface in a class

A class is the only ActionScript 3.0 language element that can implement an interface.

Use the `implements` keyword in a class declaration to implement one or more interfaces.

The following example defines two interfaces, `IAAlpha` and `IBeta`, and a class, `Alpha`, that implements them both:

```
interface IAlpha
{
    function foo(str:String):String;
}

interface IBeta
{
    function bar():void;
}

class Alpha implements IAlpha, IBeta
{
    public function foo(param:String):String {}
    public function bar():void {}
}
```

In a class that implements an interface, implemented methods must do the following:

- Use the `public` access control identifier.
- Use the same name as the interface method.
- Have the same number of parameters, each with data types that match the interface method parameter data types.
- Use the same return type.
- You do have some flexibility, however, in how you name the parameters of methods that you implement.

Although the number of parameters and the data type of each parameter in the implemented method must match that of the interface method, the parameter names do not need to match.

 Use access modifiers with classes and class members.

3.4 Under the purpose of and implement data transfer objects.

a. Using class-based models

As with MXML-based models, you cannot type the properties of a script-based model. To type properties, you must use a class-based model.

Using an ActionScript class as a model is a good option when you have to store complex data structures with typed properties, or when you want to execute client-side business logic by using application data. Also, the type information in a class-based model is retained on the server when the model is passed to a server-side data service.

The following example shows the employee model defined in an ActionScript class called `EmployeeModel`. The `EmployeeModel` model class contains various typed properties as well as accessor methods (getter/setters), a simple custom validation method, and a method that returns a Transfer Object.

Note: A Transfer Object (also known as Value Object) is a simple object that contains just the data required to carry out a certain unit of business logic. In the following example, the Employee transfer object (EmployeeTO) contains only typed properties that describe an employee.

It does not have methods like the model class. Transfer Objects are useful when moving data between the tiers by using technologies such as Flash Remoting (the transportation mechanism that is also used in Flex Data Services) and JSON that automatically serialize and deserialize complex objects across tiers while maintaining type information.

```
package
{
    public class EmployeeTO
    {
        public var firstName:String;
        public var lastName:String;
```

```

        public var department:uint;
        public var email:String;
    }
}

```



3.5 Implement accessor methods in ActionScript.

3.6 Use an ArrayCollection to sort, filter, and provide data.

a. Using an ArrayCollection

A data collection is an ordered list of data objects stored in client application memory. Flex provides an ActionScript class named `ArrayCollection` that's designed for this purpose.

More than a simple `Array`, the `ArrayCollection` class has these advantages:

Reliably executes binding expressions that refer to its stored data.

Implements a set of interfaces that provide client-side data filtering, sorting, bookmarking, and traversal.

An `ArrayCollection` can be serialized for transport over the Web in requests to Web services, remoting services, and messaging services.

a) Declaring an ArrayCollection

To declare an `ArrayCollection` in MXML, use the `<mx:ArrayCollection/>` tag and assign it an id property.

To declare an `ArrayCollection` in ActionScript, you declare and instantiate the `ArrayCollection` variable:

```

import mx.collections.ArrayCollection

private var myData:ArrayCollection;

```

b) Setting an ArrayCollection object's source property

The `ArrayCollection` class has a source property that refers to a raw `Array` containing its data. You can set an

`ArrayCollection` object's source in a number of ways:

- By passing the `Array` into the `ArrayCollection` object's constructor method.
`myData = new ArrayCollection(myArray);`

With an ActionScript statement after the `ArrayCollection` has been instantiated:

```
myData.source = ['red', 'green', 'blue'];
```

In an MXML declaration, nested in `<mx:source>` tags:

```
<mx:ArrayCollection id="myData">

<mx:source>

    <mx:String>Red</mx:String>

    <mx:String>Blue</mx:String>
    <mx:String>Green</mx:String>
</mx:source>
</mx:ArrayCollection>
```

c) Accessing data at runtime

After an `ArrayCollection` object has been created, you can dynamically get, add, and remove data at runtime with the `ArrayCollection` class interface.

The following `ArrayCollection` methods and properties are designed for this purpose:

- `addItem(item:Object)` appends a data item to the end of the collection.
- `addItemAt(item:Object, index:int)` adds a data item in the collection at the declared index position. Existing data items are shifted downward to make room for the new data item.
- `getItemAt(index:int, prefetch:int = 0)` returns a data item at the declared index position. The optional `prefetch` argument is used when an `ArrayCollection` contains managed data to indicate how many rows of data should be fetched from the server.
- `removeItemAt(index:int)` removes a data item from the `ArrayCollection` object
- `removeAll()` clears all items from the collection.
- `setItemAt(item:Object, index:int)` replaces a data item in the declared index position
- `length:int` returns the number of items in the `ArrayCollection`

d) Managing data at runtime

The `ArrayCollection` class implements a number of interfaces to allow you to dynamically manage data in client application memory at runtime. These interfaces include:

- `ICollectionView`, with methods for filtering and sorting data at runtime
- `IList`, with the methods described previously for adding, removing, and accessing data at runtime

e) Filtering Data

The `ArrayCollection` class executes filtering through its `filterFunction` property. This property is designed to reference an `ActionScript` function that you create and customize. A function designed for filtering always has this signature:

```
private function functionName( item:Object ):Boolean
```

The `item` argument can be either a generic `ActionScript` `Object` variable or a strongly typed value object. When you execute a filter, the `ArrayCollection` class loops through its source data and executes the filter function once for each data item.

If the filtering function returns `true`, the current data item is included in the resulting filtered view; if it returns `false`, the data item is hidden and won't be visible to the user unless and until the filter is removed.

The following filtering function examines a property of a data item and compares it to a value provided by the user through a visual component.

If the data item property and the user-provided value match, the function returns `true`, indicating that the data item should be included in the filtered view:

```
private function filterTables( item:Object ):Boolean
{
    return ( item.id == dg_tableData.selectedItem.id );
}
```

To use the filter function, first assign the function to the `ArrayCollection` class's filter function property by its name.

Then call the `ArrayCollection` object's `refresh()` method to cause the filtering to happen:

```
tableDataCollection.filterFunction = filterTables;
tableDataCollection.refresh();
```

f) Sorting Data

The `ArrayCollection` sorts data through use of its `sort` property. The `sort` property references an instance of the `mx.collections.Sort` class.

This class in turn has `field`'s property that references an `Array` containing instances of `mx.collections.SortField`.

The `SortField` class supports the `Boolean` properties that determine which named property of an `ArrayCollection` object's data items to sort on and how to execute the sort operation:

- `caseInsensitive` defaults to `false`, meaning that sort operations are case-sensitive by default.
- `descending` defaults to `false`, meaning that sort operations are ascending by default
- `numeric` defaults to `false`, meaning that sort operations are text-based by default

You can instantiate a `SortField` object and set all of its `Boolean` properties in the constructor method call, using this syntax

```
var mySortField:SortField = new SortField( 'propName', caseInsensitive, descending, numeric );
```

All the constructor method arguments are optional, so the following code creates a `SortField` object that sorts on a `lastName` field and uses the default settings of case-sensitive, ascending, and text:

```
var mySortField:SortField = new SortField('lastName');
```

To sort on multiple named properties, add the `SortField` objects to the array in the order of sort precedence—the first `SortField` object is the primary sort, and so on:

```
var mySort:Sort = Sort();
mySort.fields = new Array();
mySort.fields.push(new SortField('price', false, false, true));
mySort.fields.push(new SortField('title'));
```

After creating the `Sort` object and populating its `fields` property with the Array of `SortField` objects, the last step is to assign the `ArrayCollection` object's `sort` property and call its `refresh()` method:

```
myAc.sort = mySort;
myAc.refresh();
```

3.7 Implement data validation.

a. About Data Validation

The data that a user enters in a user interface might or might not be appropriate to the application. In Flex, you use a validator to ensure the values in the fields of an object meet certain criteria.

For example, you can use a validator to ensure that a user enters a valid phone number value, to ensure that a String value is longer than a set minimum length, or ensure that a ZIP code field contains the correct number of digits.

In typical client-server environments, data validation occurs on the server after data is submitted to it from the client. One advantage of using Flex validators is that they execute on the client, which lets you validate input data before transmitting it to the server.

By using Flex validators, you eliminate the need to transmit data to and receive error messages back from the server, which improves the overall responsiveness of your application.

Flex includes a set of validators for common types of user input data, including the following:

- Validating credit card numbers
- Validating currency

- Validating dates
- Validating e-mail addresses
- Validating numbers
- Validating phone numbers
- Validating using regular expressions
- Validating social security numbers
- Validating strings
- Validating ZIP codes

a) About validators

You define validators by using MXML or ActionScript. You declare a validator in MXML by using the `<mx:Validator>` tag or the tag for the appropriate validator type.

For example, to declare the standard `PhoneNumberValidator` validator, you use the `<mx:PhoneNumberValidator>` tag, as the following example shows:

```
<!-- Define the PhoneNumberValidator. -->
<mx:PhoneNumberValidator id="pnV"
    source="{phoneInput}" property="text"/>

<!-- Define the TextInput control for entering the phone number. -->
<mx:TextInput id="phoneInput"/>
<mx:TextInput id="zipCodeInput"/>
```

b) About triggering validation

You trigger validation either automatically in response to an event, or programmatically by an explicit call to the `Validator.validate()` method of a validator.

When you use events, you can cause the validator to execute automatically in response to a user action.

For example, you can use the click event of a `Button` control to trigger validation on the fields of a form, or the `valueCommit` event of a `TextInput` control to trigger validation after a user enters information in the control.

You can also trigger a validation programmatically.

For example, you might have to inspect multiple, related input fields to perform a single validation. Or you might have to perform conditional validation based on a user input.

For example, you may allow a user to select the currency used for payment, such as U.S. dollars or Euros. Therefore, you want to

make sure that you invoke a validator configured for the specified currency. In this case, you can make an explicit call to the `validate()` method to trigger the correct validator for the input value.

c) About validating required fields

Flex validators can determine when a user enters an incorrect value into a user-interface control. In addition, all validators support the required property, which, if `true`, specifies that a missing or empty value in a user-interface control causes a validation error.

The default value is `true`. Therefore, a validation error occurs by default if the user fails to enter any required data in a control associated with a validator. To disable this check, set the required property to `false`.

d) About validation errors

If a validation error occurs, by default Flex draws a red box around the component associated with the failure. If the user moves the mouse pointer over the component, Flex displays the error message associated with the error. You can customize the look of the component and the error message associated with the error.

e) About validation events

Validation is event driven. You can use events to trigger validation, programmatically create and configure validators in response to events, and listen for events dispatched by validators.

For example, when a validation operation completes, a validator dispatches a valid or invalid event, depending on the results of the validation. You can listen for these events, and then perform any additional processing that your application requires.

Alternatively, Flex components dispatch valid and invalid events, depending on the results of the validation. This lets you listen for the events being dispatched from the component being validated, rather than listening for events dispatched by the validator.

You are not required to listen for validation events. By default, Flex handles a failed validation by drawing a red box around the control that is associated with the source of the data binding.

For a successful validation, Flex clears any indicator of a previous failure

f) About custom validation

Although Flex supplies a number of predefined validators, you might find that you have to implement your own validation logic.

The `mx.validators.Validator` class is an ActionScript class that you can extend to create a subclass that encapsulates your custom validation logic.

b. Triggering validation programmatically

The `Validator` class, and all subclasses of `Validator`, include a `validate()` method that you can call to invoke a validator directly, rather than triggering the validator automatically by using an event.

The `validate()` method has the following signature:

```
validate(value:Object = null, suppressEvents:Boolean = false):ValidationResultEvent
```

The arguments have the following values:

- **value:** If `value` is `null`, use the `source` and `property` properties to specify the data to validate. If `value` is non-`null`, it specifies a field of an object relative to the `this` keyword, which means an object in the scope of the document.

You should also set the `Validator.listener` property when you specify the `value` argument. When a validation occurs, Flex applies visual changes to the object specified by the `listener` property. By default, Flex sets the `listener` property to the value of the `source` property.

However, because you do not specify the `source` property when you pass the `value` argument, you should set it explicitly

- **suppressEvents:** If `false`, dispatch either the `valid` or `invalid` event on completion. If `true`, do not dispatch events. This method returns an event object containing the results of the validation that is an instance of the `ValidationResultEvent` class.

Handling the return value of the `validate()` method

You may want to inspect the return value from the `validate()` method to perform some action when the validation succeeds or fails.

The `validate()` method returns an event object with a type defined by the `ValidationResultEvent` class. The `ValidationResultEvent` class defines several properties, including the `type` property.

The `type` property of the event object contains either `ValidationResultEvent.VALID` or `ValidationResultEvent.INVALID`, based on the validation results.

Triggering the `DateValidator` and `CreditCardValidator`

The `DateValidator` and `CreditCardValidator` can validate multiple fields by using a single validator.

A `CreditCardValidator` examines one field that contains the credit card number and a second field that contains the credit card type.

The `DateValidator` can examine a single field that contains a date, or multiple fields that together make up a date.

When you validate an object that contains multiple properties that are set independently, you often cannot use events to

automatically trigger the validator because no single field contains all of the information required to perform the validation.

One way to validate a complex object is to call the `validate()` method of the validator based on some type of user interaction.

Validating required fields

The `Validator` class, the base class for all Flex validators, contains a `required` property that when set to `true` causes validation to fail when a field is empty.

You use this property to configure the validator to fail when a user does not enter data in a required input field. You typically call the `validate()` method to invoke a validator on a required field.

This is often necessary because you cannot guarantee that an event occurs to trigger the validation— an empty input field often means that the user never gave the input control focus.

Enabling and disabling a validator

The `Validator.enabled` property lets you `enable` and `disable` a validator.

When the value of the `enabled` property is `true`, the validator is enabled; when the value is `false`, the validator is disabled.

When a validator is disabled, it dispatches no events, and the `validate()` method returns `null`.

Using data binding to configure validators

You configure validators to match your application requirements.

For example, the `StringValidator` lets you specify a minimum and maximum length of a valid string. For a `String` to be considered valid, it must be at least the minimum number of characters long, and less than or equal to the maximum number of characters.

Often, you set validator properties statically, which means that they do not change as your application executes.

For example, the following `StringValidator` defines that the input string must be at least one character long and no longer than 10 characters:

```
<mx:StringValidator required="true" minLength="1" maxLength="10"/>
```

Changing the color of the validation error message

By default, the validation error message that appears when you move the mouse pointer over a user-interface control has a red background.

You can use the `ErrorTip` style to change the color, as the following example shows:

```
<!-- Use blue for the error message. -->
```

```

<mx:Style>
    .errorTip
    {
        borderColor: #0000FF
    }
</mx:Style>

```

Showing a validation error by using `errorString`

The `UIComponent` class defines the `errorString` property that you can use to show a validation error for a component, without actually using a validator class.

When you write a `String` value to the `UIComponent.errorString` property, Flex draws a red border around the component to indicate the validation error, and the `String` appears in a `ToolTip` as the validation error message when you move the mouse over the component, just as if a validator detected a validation error.

To clear the validation error, write an empty `String`, `" "`, to the `UIComponent.errorString` property.

Note: Writing a value to the `UIComponent.errorString` property does not trigger the `valid` or `invalid` events; it only changes the border color and displays the validation error message.

Clearing a validation error

The `errorString` property is useful when you want to reset a field that is a source for validation, and prevent a validation error from occurring when you reset the field.

For example, you might provide a form to gather user input. Within your form, you might also provide a button, or other mechanism, that lets the user reset the form.

However, clearing form fields that are tied to validators could trigger a validation error.

The following example uses the `errorString` property as part of resetting the `text` property of a `TextInput` control to prevent validation errors when the form resets:

```

import mx.events.ValidationResultEvent;

private var vResult:ValidationResultEvent;

// Function to validate data and submit it to the server.
private function validateAndSubmit():void
{
    // Validate the required fields.
    vResult = zipV.validate();

```

```

        if (vResult.type==ValidationResultEvent.INVALID)
            return;
        // Submit data to server.
    }

    //Clear the input controls and the errorString property when resetting the form.
    private function resetForm():void
    {
        zipInput.text = '';
        zipInput.errorString = '';
    }

```

Specifying a listener for validation

All validators support a listener property. When a validation occurs, Flex applies visual changes to the object specified by the listener property. By default, Flex sets the listener property to the value of the source property.

That means that all visual changes that occur to reflect a validation event occur on the component being validated. However, you might want to validate one component but have validation results apply to a different component, as the following example shows:

```

<mx:ZipCodeValidator id="zipV" source="{zipCodeInput}" property="text"
listener="{errorMsg}" />
<mx:TextInput id="zipCodeInput" />
<mx:TextArea id="errorMsg" />

```

Working with validation events

- Listen for validation events dispatched by the component being validated.
- Listen for validation events dispatched by validators. Flex gives you two ways to listen for validation events: Flex components dispatch valid and invalid events, depending on the results of the validation.

This lets you listen for the events being dispatched from the component being validated, and perform any additional processing on the component based on its validation result.

The event object passed to the event listener is of type Event. All validators dispatch valid or invalid events, depending on the results of the validation. You can listen for these events, and then perform any additional processing as required by your validator.

The event object passed to the event listener is of type ValidationResultEvent.

You are not required to listen for validation events. When these events occur, by default, Flex changes the appropriate border color of the target component, displays an error message for an invalid event, or hides any previous error message for a valid event.

3.8 Manipulate XML data by using E4X.

a. About E4X

EcmaScript for XML (E4X) is an expression language that allows you to extract and manipulate data stored in XML objects with simple expressions. By using E4X, you can eliminate the significantly more verbose ActionScript code that would otherwise be required.

a) Using dot notation

An ActionScript XML object that's been created from XML notation refers to the structure's root elements, not the XML document itself.

If an HTTPService component imports the XML document and passes the ResultEvent object to an event handler function, the expression `event.result` refers to the `<invoices>` element. You typically save the returned data to a variable typed as XML and then use E4X expressions to extract data as needed:

```
[Bindable] private var xInvoices:XML;

private function resultHandler( event:ResultEvent ):void
{
    xInvoices = event.result as XML;
}
```

Starting at the root element, you then “walk” down the XML hierarchy one level at a time using simple dot notation and element names. So for example, this code extracts all elements named `<invoice>` that are child elements of the root:

```
xResult = xInvoices.invoice;
```

The returned XMLList looks like this:

```
<invoice>
```

```
    <customer>
        <firstname>Jonnie</firstname>
        <lastname>Spratley</lastname>
    </customer>
    <items>
        ...
    </items>
```

```
</invoice>
```

b) Using array notation

When two or more elements have the same name at the same level of the XML hierarchy, you can refer to individual items as Array elements. As with all ActionScript array and index notation, indexing starts at 0. This statement extracts the second element named invoice:

```
xReturn = xInvoices.invoice[1];
```

The resulting XML object looks like this:

```
<invoice>
```

```
    <customer>
        <firstname>Jonnie</firstname>
        <lastname>Spratley</lastname>
    </customer>
    <items>
        <lineitem price="4.00" quantity="3">Mouse Pad</lineitem>
        <lineitem price="6.00" quantity="1">USB Adapter</lineitem>
    </items>
</invoice>
```

After you've found an element using array notation, you can continue down the XML hierarchy with extended dot notation. This code extracts the customer element that's a child of the first invoice element:

```
xReturn = xInvoices.invoice[0].customer
```

The resulting XML object looks like this:

```
<customer>
    <firstname>Jonnie</firstname>
    <lastname>Spratley</lastname>
</customer>
```

c) Using the descendant accessor operator

A single dot (.) causes the expression to look at direct child elements of the current XML object. You also can do a deep search of the XML hierarchy for elements by their names with the double-dot (..) descendant accessor operator. This allows you to search for all elements of the same name, regardless of their position or how many parent elements there are between the current element and the one you're looking for.

This code extracts all <customer> elements regardless of their position in the content:

```
xReturn = xInvoices..customer;
```

The resulting XMLList object looks like this:

```
<customer>
```

```

        <firstname>Jonnie</firstname>
        <lastname>Spratley</lastname>
    </customer>
    <customer>
        <firstname>Foo</firstname>
        <lastname>Bar</lastname>
    </customer>

```

d) Filtering XML data with predicate expressions

A predicate expression lets you filter data in an XML object, accomplishing tasks similar to the WHERE clause in an SQL statement. The predicate expression itself is an ActionScript comparison expression wrapped in parentheses. You appended the predicate to the part of the E4X expression that indicates what data you want to return, separated with a dot operator.

This ActionScript extracts all <customer> elements that have a <lastname> matching the String value Jones:

```
xRetrun = xInvoices..customer.(lastname=='Jones')
```

The resulting XML node looks like this:

```

<customer>

    <firstname>Foo</firstname>
    <lastname>Jones</lastname>
</customer>

```

4. Interacting with data sources and servers

4.1 Implement simple LiveCycle Data Services (LCDS) messaging and data management.

a. About LiveCycle Data Services

LiveCycle Data Services provides a messaging infrastructure for building data-rich Flex applications. This infrastructure underlies the LiveCycle Data Services features that are designed for moving data to and from applications: RPC services, the Data Management Service, and the Message Service.

The following table describes the types of services provided in LiveCycle Data Services:

Service	Description
RPC services	Provides a call and response model for accessing external data. Lets you create applications that make asynchronous requests to remote services that process the requests, and then return data to your Flex application. For more information, see About RPC services.

Data Management Service	<p>Provides data synchronization between application tiers, real-time data updates, on-demand data paging, data replication, occasionally connected application services, and integration with data sources through adapters. Lets you create applications that work with distributed data, and lets you manage large collections of data and nested data relationships, such as one-to-one and one-to-many relationships.</p> <p>For more information, see About data management.</p>
Message Service	<p>Provides messaging services for collaborative and real-time applications. Lets you create applications that can send messages to and receive messages from other applications, including Flex applications and Java Message Service (JMS) applications.</p> <p>For more information, see About messaging.</p>

a) About data management

The Data Management Service is a part of LiveCycle Data Services that lets you create applications that work with distributed data. By using the Data Management Service, you build applications that provide data synchronization, data replication, on-demand data paging, and occasionally connected application services. Additionally, you can manage large collections of data and nested data relationships, such as one-to-one and one-to-many relationships, and use data adapters to integrate with data resources.

A client-side DataService component, which you can create in MXML or ActionScript, calls methods on a server-side Data Management Service destination to perform such activities as filling client-side data collections with data from remote data sources and synchronizing the client and server versions of data.

b) About messaging

The messaging capability in LiveCycle Data Services is built on the same underlying messaging infrastructure as the RPC services and Data Management Service. By using messaging components, you create collaborative and real-time messaging applications. These components provide a client-side API for passing text and objects to and from messaging clients, and server-side configuration of message destinations.

You can create messaging components in MXML or ActionScript that connect to server-side message destinations, send messages to those destinations, and receive messages from other messaging clients. Components that send messages are called producers, and those that receive messages are called consumers.

Messages are sent over transport channels, which use specific transport protocols, such as the Realtime Message Protocol (RTMP) and Action Message Format (AMF). Messaging clients can be Flex applications or other types of applications, such as Java Message Service (JMS) applications. JMS is a Java API that lets applications create, send, receive, and read messages. JMS applications can

publish and subscribe to the same message destinations as Flex applications. However, you can create a wide variety of messaging applications using just Flex messaging.

LCDS messaging could originate from the following places:

- Java Message Service (JMS)
- ColdFusion event gateways
- Other Flex clients using the `<mx:Producer>` tag

Important terms to understand when working with Flex Messaging:

- **Message:** A piece of data sent from a message producer through the messaging service to a consumer. A message consists of a header that describes the message and a body that contains the data for the message.
- **Consumer:** An application that makes use of messages. Dashboard applications that monitor the sales of a site would be considered a consumer, because you are subscribing to that data.
- **Producer:** The sender of a message.
- **Destination:** An identifier for a place to which a producer sends a message. The LCDS will route the message appropriately, based on its destination. This architecture makes it possible for the producers and consumers to be entirely decoupled, knowing nothing about each other.
- **Channel:** A gateway to a message destination. Producers use channels to connect to the destination within Flex Data Services, and FDS uses channels to connect to the listening consumers.
- **Adapter:** Elements that translate message between the FDS and external messaging services. LCDS ships with adapters for JMS, ColdFusion Gateways, and the ActionScript adapter, which is used to send messages to the Flash Clients.

c) Using messaging in a Flex application

A Flex client application uses the client-side messaging API to send messages to, and receive messages from, a server-side destination. Messages are sent to and from destinations over a protocol-specific message channel.

The two primary client-side messaging components are Producer and Consumer components. A Producer component sends messages to a server-side destination. A Consumer component subscribes to a server-side destination and receives messages that a Producer component sends to that destination. You can create Producer and Consumer components in MXML or ActionScript.

Producer and Consumer components both require a valid message destination that you configure in the Flex services configuration file. For information about message destinations, see [Configuring Messaging on the Server](#).

A Flex application often contains at least one pair of Producer and Consumer components. This enables each application to send messages to a destination and receive messages that other applications send to that destination. It is very easy to test this type of application by running it in different browser windows and sending messages from each of them.

d) Working with Producer components

You can create Producer components in MXML or ActionScript. To send a message, you create an `AsyncMessage` object and then call a Producer component's `send()` method to send the message to a destination.

You can also specify acknowledge and fault event handlers for a Producer component. An acknowledge event is broadcast when a destination successfully receives a message that a Producer component sends. A fault event is dispatched when a destination cannot successfully process a message due to a connection-, server-, or application-level failure.

A Producer component can send messages to destinations that use the Flex Message Service with no message adapters or to destinations that use message adapters. A message adapter is server-side code that acts as a conduit between the Flex Message Service and other messaging systems.

For example, the Java Message Service (JMS) adapter enables Producer components to send messages to JMS topics. On the client side, the APIs for non-adapter and adapter-enabled destinations are identical. The definition for a particular destination in the Flex services configuration file determines what happens when a Producer component sends messages to the destination or a Consumer component subscribes to the destination.

e) Creating a Producer component in MXML

You use the `<mx:Producer>` tag to create a Producer component in MXML. The tag must contain an `id` value. It should also specify a destination that is defined in the server-side `services-config.xml` file.

The following code shows an `<mx:Producer>` tag that specifies a destination and acknowledge and fault event handlers:

```
<mx:Application xmlns:mx="http://www.adobe.com/2006/mxml">
  <mx:Script>
    <![CDATA[
      import mx.rpc.events.FaultEvent;
      import mx.messaging.*;
      import mx.messaging.messages.*;
      import mx.messaging.events.*;
      private function messageHandler(event:MessageEvent):void {
        // Handle message event.
      }
      private function acknowledgeHandler(event:MessageAckEvent):void{
        // Handle message event.
      }
      private function faultHandler(event:MessageFaultEvent):void {
        // Handle fault event.
      }
    ]]>
  </mx:Script>
  <mx:Producer id="producer"
    destination="ChatTopicJMS"
    acknowledge="acknowledgeHandler(event)"
    fault="faultHandler(event)"/>
</mx:Application>
```

f) Creating a Producer component in ActionScript

You can create a Producer component in an ActionScript method. The following code shows a Producer component that is created in a method in an `<mx:Script>` tag. The import statements import the classes required to create a Producer component, create Message objects, add event listeners, and create message handlers.

```
<mx:Script>
    <![CDATA [
        import mx.messaging.*;
        import mx.messaging.messages.*;
        import mx.messaging.events.*;
        private var producer:Producer;
        private function acknowledgeHandler(event:MessageAckEvent):void{
            // Handle message event.
        }
        private function faultHandler(event:MessageFaultEvent):void{
            // Handle fault event.
        }
        private function logon():void {
            producer = new Producer();
            producer.destination = "ChatTopicJMS";
            producer.addEventListener(MessageAckEvent.ACKNOWLEDGE, acknowledgeHandler);
            producer.addEventListener(MessageFaultEvent.FAULT, faultHandler);
        }
    ]]>
</mx:Script>
```

g) Sending a message to a destination

To send a message from a Producer component to a destination, you create an `mx.messaging.messages.AsyncMessage` object, populate the body of the `AsyncMessage` object, and then call the component's `send()` method. You can create text messages and messages that contain objects.

The following code creates a message, populates the body of the message with text, and sends the message by calling a Producer component's `send()` method:

```
<mx:Application xmlns:mx="http://www.adobe.com/2006/mxml">
    <mx:Script>
        <![CDATA [
            import mx.messaging.*;
            import mx.messaging.messages.*;
            import mx.messaging.events.*;
            private var producer:Producer;

            private function sendMessage():void {
                var message:AsyncMessage = new AsyncMessage();
                message.body = userName.text + ": " + input.text;
                producer.send(message);
            }
        ]]>
    </mx:Script>
    <mx:TextInput id="userName"/>
    <mx:TextInput id="input"/>
    <mx:Button label="Send" click="sendMessage();" />
</mx:Application>
```

The following code creates a message, populates the body of the message with an object, and sends the message by calling a Producer component's `send()` method.

```
<mx:Application xmlns:mx="http://www.adobe.com/2006/mxml">
  <mx:Script>
    <![CDATA[
      import mx.messaging.*;
      import mx.messaging.messages.*;
      import mx.messaging.events.*;
      private var producer:Producer;

      private function sendMessage():void {
        var message:AsyncMessage = new AsyncMessage();
        message.body=new Object();
        message.body.prop1 = "abc";
        message.body.prop2 = "abc";
        message.body.theCollection=['b','a',3,new Date()];
        producer.send(message);
      }
    ]]>
  </mx:Script>
</mx:Application>
```

h) Adding extra information to a message

You can include extra information for a message in the form of message headers. You can send strings and numbers in message headers. The headers of a message are contained in an associative array where the key is the header name. The headers property of a message class lets you add headers for a specific message instance.

The following code adds a message header called `prop1` and sets its value:

```
<mx:Application xmlns:mx="http://www.adobe.com/2006/mxml">
  <mx:Script>
    <![CDATA[
      import mx.messaging.*;
      import mx.messaging.messages.*;
      import mx.messaging.events.*;
      private var producer:Producer;

      private function sendMessage():void {
        var message:AsyncMessage = new AsyncMessage();
        message.headers = new Array();
        message.headers["prop1"] = 5;
        message.body =input.text;
        producer.send(message);
      }
    ]]>
  </mx:Script>
  <mx:TextInput id="input"/>
</mx:Application>
```

You can use a Consumer component's selector property to filter the messages that the component receives, based on message

header values.

i) Working with Consumer components

You can create Consumer components in MXML or ActionScript. To subscribe to a destination, you call a Consumer component's `subscribe()` method.

You can also specify message and fault event handlers for a Consumer component. A message event is broadcast when a message is received by the destination to which a Consumer component is subscribed.

A fault event is broadcast when the channel to which the Consumer component is subscribed can't establish a connection to the destination, the subscription request is denied, or if a failure occurs when the Consumer component's `receive()` method is called. The `receive()` method retrieves any pending messages from the destination.

j) Creating a Consumer component in MXML

You use the `<mx:Consumer>` tag to create a Consumer component in MXML. The tag must contain an `id` value. It should also specify a destination that is defined in the server-side `services-config.xml` file.

The following code shows an `<mx:Consumer>` tag that specifies a destination and acknowledge and fault event handlers:

```
<mx:Application xmlns:mx="http://www.adobe.com/2006/mxml">
  <mx:Script>
    <![CDATA[
      import mx.messaging.*;
      import mx.messaging.messages.*;
      import mx.messaging.events.*;
      private function messageHandler(event:MessageEvent):void {
        // Handle message event.
      }
      private function faultHandler(event:MessageFaultEvent):void {
        // Handle fault event.
      }
    ]]>
  </mx:Script>
  <mx:Consumer id="consumer"
    destination="ChatTopicJMS"
    message="messageHandler(event)"
    fault="faultHandler(event)"/>
</mx:Application>
```

k) Creating a Consumer component in ActionScript

You can create a Consumer component in an ActionScript method. The following code shows a Consumer component created in a method in an `<mx:Script>` tag. The import statements import the classes required to create a Consumer component, create `AsyncMessage` objects, add event listeners, and create message handlers.

```
<mx:Application xmlns:mx="http://www.adobe.com/2006/mxml">
  <mx:Script>
    <![CDATA[
```

```

import mx.messaging.*;
import mx.messaging.messages.*;
import mx.messaging.events.*;
private var consumer:Consumer;
private function acknowledgeHandler(event:MessageAckEvent):void{
    // Handle message event.
}
private function faultHandler(event:MessageFaultEvent):void{
    // Handle fault event.
}
private function logon():void {
    consumer = new Consumer();
    consumer.destination = "ChatTopicJMS";
    consumer.addEventListener
        (MessageAckEvent.ACKNOWLEDGE, acknowledgeHandler);
    consumer.addEventListener
        (MessageFaultEvent.FAULT, faultHandler);
}
]]>
</mx:Script>
</mx:Application>

```

I) Subscribing to a destination

Whether you create a Consumer component in MXML or ActionScript, you still must call the component's `subscribe()` method to subscribe to a destination and receive messages from that destination.

A Consumer component can subscribe to destinations that use the LiveCycle Data Services Message Service with no message adapters or to destinations that use message adapters. For example, the Java Message Service (JMS) adapter enables Consumer components to subscribe to JMS topics. On the client side, the API for non-adapter and adapter-enabled destinations is identical. The definition for a particular destination in the services configuration file determines what happens when a Consumer component subscribes to the destination or a Producer component sends messages to the destination. For information about configuring destinations, see [Configuring Messaging on the Server](#).

The following code shows a call to a Consumer component's `subscribe()` method:

```

<mx:Script>
<![CDATA[
    import mx.messaging.*;
    import mx.messaging.messages.*;
    import mx.messaging.events.*;
    private var consumer:Consumer;

    private function logon():void {
        consumer = new Consumer();
        consumer.subscribe();
    }
]]>
</mx:Script>

```

4.1 - Create, connect to, and define a local database.

a. Creating a database

To create a new database file, you first create an `SQLConnection` instance. You call its `open()` method to open it in synchronous execution mode, or its `openAsync()` method to open it in asynchronous execution mode.

The `open()` and `openAsync()` methods are used to open a connection to a database. If you pass a `File` instance that refers to a non-existent file location for the reference parameter (the first parameter), the `open()` or `openAsync()` method will create a database file at that file location and open a connection to the newly created database.

Whether you call the `open()` method or the `openAsync()` method to create a new database, the database file's name can be any valid file name, with any file extension. If you call the `open()` or `openAsync()` method with null for the reference parameter, a new in-memory database is created rather than a database file on disk.

The following code listing shows the process of creating a database file (a new database) using asynchronous execution mode. In this case, the database file is saved in the application's storage directory, with the file name "DBSample.db":

```
import flash.data.SQLConnection;
import flash.events.SQLErrorEvent;
import flash.events.SQLEvent;
import flash.filesystem.File;

var conn:SQLConnection = new SQLConnection();
conn.addEventListener(SQLEvent.OPEN, openHandler);
conn.addEventListener(SQLErrorEvent.ERROR, errorHandler);

var dbFile:File = File.applicationStorageDirectory.resolvePath("DBSample.db");

conn.openAsync(dbFile);

function openHandler(event:SQLEvent):void
{
    trace("the database was created successfully");
}

function errorHandler(event:SQLErrorEvent):void
{
    trace("Error message:", event.error.message);
    trace("Details:", event.error.details);
}
```

To execute operations synchronously, when you open a database connection with the `SQLConnection` instance, call the `open()` method.

The following example shows how to create and open an `SQLConnection` instance that executes its operations synchronously:

```
import flash.data.SQLConnection;
import flash.events.SQLErrorEvent;
```



```

import flash.events.SQLEvent;
import flash.filesystem.File;

var conn:SQLConnection = new SQLConnection();

var dbFile:File = File.applicationStorageDirectory.resolvePath("DBSample.db");

try
{
    conn.open(dbFile);
    trace("the database was created successfully");
}
catch (error:SQLError)
{
    trace("Error message:", error.message);
    trace("Details:", error.details);
}

```

a) Creating database tables

Creating a table in a database involves executing a SQL statement on that database, using the same process that you use to execute a SQL statement such as SELECT, INSERT, and so forth. To create a table, you use a CREATE TABLE statement, which includes definitions of columns and constraints for the new table. For more information about executing SQL statements, see Working with SQL statements.

The following example demonstrates creating a new table named "employees" in an existing database file, using asynchronous execution mode. Note that this code assumes there is a SQLConnection instance named conn that is already instantiated and is already connected to a database.

```

import flash.data.SQLConnection;
import flash.data.SQLStatement;
import flash.events.SQLErrorEvent;
import flash.events.SQLEvent;

// ... create and open the SQLConnection instance named conn ...

var createStmt:SQLStatement = new SQLStatement();
createStmt.sqlConnection = conn;

var sql:String =
    "CREATE TABLE IF NOT EXISTS employees (" +
    "    empId INTEGER PRIMARY KEY AUTOINCREMENT, " +
    "    firstName TEXT, " +
    "    lastName TEXT, " +
    "    salary NUMERIC CHECK (salary > 0) " +
    ")";
createStmt.text = sql;

createStmt.addEventListener(SQLEvent.RESULT, createResult);
createStmt.addEventListener(SQLErrorEvent.ERROR, createError);

```

```

createStmt.execute();

function createResult(event:SQLEvent):void
{
    trace("Table created");
}

function createError(event:SQLErrorEvent):void
{
    trace("Error message:", event.error.message);
    trace("Details:", event.error.details);
}

```

The following example demonstrates how to create a new table named "employees" in an existing database file, using synchronous execution mode. Note that this code assumes there is a `SQLConnection` instance named `conn` that is already instantiated and is already connected to a database.

```

import flash.data.SQLConnection;
import flash.data.SQLStatement;
import flash.events.SQLErrorEvent;
import flash.events.SQLEvent;

// ... create and open the SQLConnection instance named conn ...

var createStmt:SQLStatement = new SQLStatement();
createStmt.sqlConnection = conn;

var sql:String =
    "CREATE TABLE IF NOT EXISTS employees (" +
    "    empId INTEGER PRIMARY KEY AUTOINCREMENT, " +
    "    firstName TEXT, " +
    "    lastName TEXT, " +
    "    salary NUMERIC CHECK (salary > 0)" +
    ")";
createStmt.text = sql;

try
{
    createStmt.execute();
    trace("Table created");
}
catch (error:SQLError)
{
    trace("Error message:", error.message);
    trace("Details:", error.details);
}

```

b) Connecting to a database

Before you can perform any database operations, first open a connection to the database file. An `SQLConnection` instance is used to represent a connection to one or more databases. The first database that is connected using an `SQLConnection` instance is known as the "main" database.

* This database is connected using the `open()` method (for synchronous execution mode) or the `openAsync()` method (for asynchronous execution mode).

If you open a database using the asynchronous `openAsync()` operation, register for the `SQLConnection` instance's open event in order to know when the `openAsync()` operation completes. Register for the `SQLConnection` instance's error event to determine if the operation fails.

The following example shows how to open an existing database file for asynchronous execution. The database file is named "DBSample.db" and is located in the user's application storage directory.

```
import flash.data.SQLConnection;
import flash.data.SQLMode;
import flash.events.SQLErrorEvent;
import flash.events.SQLEvent;
import flash.filesystem.File;

var conn:SQLConnection = new SQLConnection();
conn.addEventListener(SQLEvent.OPEN, openHandler);
conn.addEventListener(SQLErrorEvent.ERROR, errorHandler);

var dbFile:File = File.applicationStorageDirectory.resolvePath("DBSample.db");

conn.openAsync(dbFile, SQLMode.UPDATE);

function openHandler(event:SQLEvent):void
{
    trace("the database opened successfully");
}

function errorHandler(event:SQLErrorEvent):void
{
    trace("Error message:", event.error.message);
    trace("Details:", event.error.details);
}
```

The following example shows how to open an existing database file for synchronous execution. The database file is named "DBSample.db" and is located in the user's application storage directory.

```

import flash.data.SQLConnection;
import flash.data.SQLMode;
import flash.events.SQLErrorEvent;
import flash.events.SQLEvent;
import flash.filesystem.File;

var conn:SQLConnection = new SQLConnection();

var dbFile:File = File.applicationStorageDirectory.resolvePath("DBSample.db");

try
{
    conn.open(dbFile, SQLMode.UPDATE);
    trace("the database opened successfully");
}
catch (error:SQLError)
{
    trace("Error message:", error.message);
    trace("Details:", error.details);
}

```

Notice that in the `openAsync()` method call in the asynchronous example, and the `open()` method call in the synchronous example, the second argument is the constant `SQLMode.UPDATE`. Specifying `SQLMode.UPDATE` for the second parameter (`openMode`) causes the runtime to dispatch an error if the specified file doesn't exist.

If you pass `SQLMode.CREATE` for the `openMode` parameter (or if you leave the `openMode` parameter off), the runtime attempts to create a database file if the specified file doesn't exist.

You can also specify `SQLMode.READ` for the `openMode` parameter to open an existing database in a read-only mode. In that case data can be retrieved from the database but no data can be added, deleted, or changed.

4.3 - Add, update, and remove records from local database.

a. Inserting Data

a) Executing an INSERT statement

To add data to a table in a database, you create and execute an `SQLStatement` instance whose text is a SQL INSERT statement.

The following example uses an `SQLStatement` instance to add a row of data to the already-existing employees table. This example demonstrates inserting data using asynchronous execution mode.

Note that this listing assumes that there is an `SQLConnection` instance named `conn` that has already been instantiated and is already connected to a database. It also assumes that the "employees" table has already been created.

```

import flash.data.SQLConnection;
import flash.data.ResultSet;
import flash.data.SQLStatement;
import flash.events.SQLErrorEvent;

```

```

import flash.events.SQLEvent;

// ... create and open the SQLConnection instance named conn ...

// create the SQL statement
var insertStmt:SQLStatement = new SQLStatement();
insertStmt.sqlConnection = conn;

// define the SQL text
var sql:String =
    "INSERT INTO employees (firstName, lastName, salary) " +
    "VALUES ('Bob', 'Smith', 8000)";
insertStmt.text = sql;

// register listeners for the result and failure (status) events
insertStmt.addEventListener(SQLEvent.RESULT, insertResult);
insertStmt.addEventListener(SQLErrorEvent.ERROR, insertError);

// execute the statement
insertStmt.execute();

function insertResult(event:SQLEvent):void
{
    trace("INSERT statement succeeded");
}

function insertError(event:SQLErrorEvent):void
{
    trace("Error message:", event.error.message);
    trace("Details:", event.error.details);
}

```

b) Changing or deleting data

The process for executing other data manipulation operations is identical to the process used to execute a SQL SELECT or INSERT statement. Simply substitute a different SQL statement in the SQLStatement instance's text property:

- To change existing data in a table, use an UPDATE statement.
- To delete one or more rows of data from a table, use a DELETE statement.

4.5 - Interact with remote data and services by using Remote Procedure Call (RPC) services.

a. About RPC Services

a) Calling HTTP services in ActionScript

The following example shows an HTTP service call in an ActionScript script block. Calling the `useHTTPService()` method declares the service, sets the destination, sets up result and fault event listeners, and calls the service's `send()` method.

```
private var service:HTTPService

public function useHttpService(parameters:Object):void
{
    service = new HTTPService();
    service.destination = "sampleDestination";
    service.method = "POST";
    service.addEventListener("result", httpResult);
    service.addEventListener("fault", httpFault);
    service.send(parameters);
}

public function httpResult(event:ResultEvent):void
{
    var result:Object = event.result;
    //Do something with the result.
}

public function httpFault(event:FaultEvent):void
{
    var faultstring:String = event.fault.faultString;
    Alert.show(faultstring);
}
```

b) Calling web services in ActionScript

The following example shows a web service call in an ActionScript script block. Calling the `useWebService()` method declares the service, sets the destination, fetches the WSDL document, and calls the `echoArgs()` method of the service.

```
import mx.rpc.soap.WebService;
import mx.rpc.events.ResultEvent;
```

```

import mx.rpc.events.FaultEvent;

public function useWebService(intArg:int, strArg:String):void
{
    var ws:WebService = new WebService();
    ws.destination = "echoArgService";
    ws.echoArgs.addEventListener("result", echoResultHandler);
    ws.addEventListener("fault", faultHandler);
    ws.loadWSDL();
    ws.echoArgs(intArg, strArg);
}

public function echoResultHandler(event:ResultEvent):void
{
    var retStr:String = event.result.echoStr;
    var retInt:int = event.result.echoInt;
    //Do something.
}

public function faultHandler(event:FaultEvent):void
{
    //deal with event.fault.faultString, etc
}

```

c) Calling RemoteObject components in ActionScript

In the following ActionScript example, calling the `useRemoteObject()` method declares the service, sets the destination, sets up result and fault event listeners, and calls the service's `getList()` method.

```

import mx.controls.Alert;
import mx.rpc.remoting.RemoteObject;
import mx.rpc.events.ResultEvent;
import mx.rpc.events.FaultEvent;

[Bindable]
public var empList:Object;
public var employeeRO:RemoteObject;

public function useRemoteObject(intArg:int, strArg:String):void
{
    employeeRO = new RemoteObject();
    employeeRO.destination = "SalaryManager";
    employeeRO.getList.addEventListener("result", getListResultHandler);
    employeeRO.addEventListener("fault", faultHandler);
    employeeRO.getList(deptComboBox.selectedItem.data);
}

public function getListResultHandler(event:ResultEvent):void
{
    // Do something
    empList = event.result;
}

```

```

}

public function faultHandler (event:FaultEvent):void
{
    // Deal with event.fault.faultString, etc.
}

```

4.6 Upload files to a server.

a. The following is working with file upload.

a) Working with file upload

The `FileReference` class lets you add the ability to upload and download files between a client and a server. Users are prompted to select a file to upload or a location for download from a dialog box (such as the Open dialog box on the Windows operating system).

Each `FileReference` object that you create with ActionScript refers to a single file on the user's hard disk. The object has properties that contain information about the file's size, type, name, creation date, and modification date.

Note: The creator property is supported on Mac OS only. All other platforms return null.

You can create an instance of the `FileReference` class in two ways. You can use the new operator, as the following code shows:

```

import flash.net.FileReference;

private var myFileReference:FileReference = new FileReference();

```

Or you can call the `FileReferenceList.browse()` method, which opens a dialog box on the user's system to prompt the user to select one or more files to upload and then creates an array of `FileReference` objects if the user selects one or more files successfully.

Each `FileReference` object represents a file selected by the user from the dialog box. A `FileReference` object does not contain any data in the `FileReference` properties (such as name, size, or modificationDate) until one of the following happens:

The `FileReference.browse()` method or `FileReferenceList.browse()` method has been called, and the user has selected a file from the file picker.

The `FileReference.download()` method has been called, and the user has selected a file from the file picker.

Note: When performing a download, only the `FileReference.name` property is populated before the download is complete. After the file has been downloaded, all properties are available.

While calls to the `FileReference.browse()`, `FileReferenceList.browse()`, or `FileReference.download()` method are executing, most players will continue SWF file playback.

b) FileReference class

The `FileReference` class allows you to upload and download files between a user's computer and a server. An operating system dialog box prompts the user to select a file to upload or a location for download.

Each `FileReference` object refers to a single file on the user's disk and has properties that contain information about the file's size, type, name, creation date, modification date, and creator.

`FileReference` instances can be created in two ways:

When you use the new operator with the `FileReference` constructor, as in the following:

```
var myFileReference:FileReference = new FileReference();
```

When you call `FileReferenceList.browse()`, which creates an array of `FileReference` objects.

For uploading and downloading operations, a SWF file can access files only within its own domain, including any domains specified by a cross-domain policy file.

You need to put a policy file on the file server if the SWF file initiating the upload or download doesn't come from the same domain as the file server.

Note: You can only perform one `browse()` or `download()` action at a time, because only one dialog box can be open at any point.

The server script that handles the file upload should expect an HTTP POST request with the following elements:

Content-Type with a value of `multipart/form-data`.

Content-Disposition with a name attribute set to "Filedata" and a filename attribute set to the name of the original file.

You can specify a custom name attribute by passing a value for the `uploadDataFieldName` parameter in the `FileReference.upload()` method.

c) Uploading files to a server

To upload files to a server, first call the `browse()` method to allow a user to select one or more files.

Next, when the `FileReference.upload()` method is called, the selected file will be transferred to the server. If the user selected multiple files using the `FileReferenceList.browse()` method, Flash Player creates an array of selected files called `FileReferenceList.fileList`. You can then use the `FileReference.upload()` method to upload each file individually.

Note: Using the `FileReference.browse()` method allows you to upload single files only. To allow a user to upload multiple files, you must use the `FileReferenceList.browse()` method.

By default, the system file picker dialog box allows users to pick any file type from the local computer, although developers can specify one or more custom file type filters by using the `FileFilter` class and passing an array of file filter instances to the `browse()` method:

```
var imageTypes:FileFilter = new FileFilter( "Images (*.jpg, *.png)", "*.jpg; *.png" );
var textTypes:FileFilter = new FileFilter( "Text Files (*.txt, *.rtf)", "*.txt; *.rtf" );
var allTypes:Array = new Array( imageTypes, textTypes );
var fileRef:FileReference = new FileReference();
fileRef.browse(allTypes);
```

When the user has selected the files and clicked the Open button in the system file picker, the `Event.SELECT` event is dispatched. If the `FileReference.browse()` method was used to select a file to upload, the following code is needed to send the file to a web server:

```
private function uploadFile():void
{
    var fileRef:FileReference = new FileReference();
    fileRef.addEventListener(Event.SELECT, selectHandler);
    fileRef.addEventListener(Event.COMPLETE, completeHandler);

    try {
        var success:Boolean = fileRef.browse();
    } catch (error:Error) {
        trace("Unable to browse for files.");
    }
}

private function selectHandler(event:Event):void
{
    var request:URLRequest = new URLRequest("http://yourdomain.com/fileupload.cfm")
    try {
        fileRef.upload(request);
    } catch (error:Error) {
        trace("Unable to upload file.");
    }
}

private function completeHandler(event:Event):void
{
    trace("uploaded");
}
```

You can send data to the server with the `FileReference.upload()` method by using the `URLRequest.method` and `URLRequest.data` properties to send variables using the `POST` or `GET` methods.

When you attempt to upload a file using the `FileReference.upload()` method, any of the following events may be dispatched:

`Event.OPEN`: Dispatched when an upload operation starts.

`ProgressEvent.PROGRESS`: Dispatched periodically during the file upload operation.

`Event.COMPLETE`: Dispatched when the file upload operation completes successfully.

`SecurityErrorEvent.SECURITY_ERROR`: Dispatched when an upload fails because of a security violation.

`HTTPStatusEvent.HTTP_STATUS`: Dispatched when an upload fails because of an HTTP error.

`IOErrorEvent.IO_ERROR`: Dispatched if the upload fails because of any of the following reasons:

An input/output error occurred while Flash Player is reading, writing, or transmitting the file.

The SWF tried to upload a file to a server that requires authentication (such as a user name and password). During upload, Flash Player does not provide a means for users to enter passwords.

The `url` parameter contains an invalid protocol. The `FileReference.upload()` method must use either `HTTP` or `HTTPS`.

You can pass additional variables to the upload script using either the `POST` or `GET` request method. To send additional `POST` variables to your upload script, you can use the following code:

```
var fileRef:FileReference = new FileReference();
fileRef.addEventListener(Event.SELECT, selectHandler);
fileRef.addEventListener(Event.COMPLETE, completeHandler);
fileRef.browse();
private function selectHandler(event:Event):void
{
    var params:URLVariables = new URLVariables();
    params.date = new Date();
    params.ssid = "94103-1394-2345";
    var request:URLRequest = new URLRequest("http://yourdomain.com/fileupload.cfm");
    request.method = URLRequestMethod.POST;
    request.data = params;

    fileRef.upload(request, "Custom1");
}

private function completeHandler(event:Event):void
{
    trace("uploaded");
}
```

The previous example creates a new `URLVariables` object that you pass to the remote server-side script.

In previous versions of `ActionScript`, you could pass variables to the server upload script by passing values in the query string. `ActionScript 3.0` allows you to pass variables to the remote script using a `URLRequest` object, which allows you to pass data using either the `POST` or `GET` method; this, in turn, makes passing larger sets of data easier and cleaner.

In order to specify whether the variables are passed using the `GET` or `POST` request method, you can set the `URLRequest.method` property to either `URLRequestMethod.GET` or `URLRequestMethod.POST`, respectively.

`ActionScript 3.0` also lets you override the default `Filedata` upload file field name by providing a second parameter to the `upload()` method, as demonstrated in the previous example (which replaced the default value `Filedata` with `Custom1`).

By default, Flash Player will not attempt to send a test upload, although you can override this by passing a value of `true` as the third parameter to the `upload()` method.

The purpose of the test upload is to check whether the actual file upload will be successful and that server authentication, if required, will succeed.

5. Using Flex in the Adobe Integrated Runtime (AIR)

5.1 - Given a scenario, compile and export a release build of an AIR application.

a. Creating a release build of your application in Flex Builder

When you create and run an application in Flex Builder, the Flex compiler includes debug information in that application so that you can set breakpoints and view variables and perform other debugging tasks. However, the SWF file that Flex Builder generates by default includes debugging information that makes it larger than the release build of the SWF file.

To create a release build of your application in Flex Builder, select `Project > Export Release Build`. This compiles a version of your application's SWF file that does not contain any debug information in it. This SWF file is smaller than the SWF files you compile normally. This also compiles any modules that are in the application's project without debug information.

In general, all compiled SWF files that are stored in the project's `/bin-debug` directory contain debug information. When you export the application, you choose a new output directory. The default is the `/bin-release` directory.

To compile a release build on the command line, set the debug compiler option to `false`. This prevents debug information from being included in the final SWF file.

a) Enabling accessibility

The Flex accessibility option lets you create applications that are accessible to users with disabilities. By default, accessibility is disabled.

You enable the accessibility features of Flex components at compile time by using options to a Flex Builder project, setting the accessible option to true for the command-line compiler, or setting the <accessible> tag in the flex-config.xml file to true.

b) Preventing users from viewing your source code

Flex lets you publish your source code with your deployed application. You might want to enable this option during the development process, but disable it for deployment. Or, you might want to include your source code along with your deployed application.

In Flex Builder, the Export Release Build wizard lets you specify whether to publish your source code.

You can also use the viewSourceURL property of the Application class to set the URL of your source code.

5.2 - Create, populate, and delete files and directories on a local file system.

Adobe AIR provides classes that you can use to access, create, and manage both files and folders. These classes,

Classes contained in the flash.filesystem package, are used as follows:

File Classes	Description
File	File object represents a path to a file or directory. You use a file object to create a pointer to a file or folder, initiating interaction with the file or folder.
FileMode	<p>The FileMode class defines string constants used in the fileMode parameter of the open() and openAsync() methods of the FileStream class.</p> <p>The fileMode parameter of these methods determines the capabilities available to the FileStream object once the file is opened, which include writing, reading, appending, and updating.</p>
FileStream	FileStream object is used to open files for reading and writing. Once you've created a File object that points to a new or existing file, you pass that pointer to the FileStream object so that you can open and then manipulate data within the file.

Some methods in the File class have both synchronous and asynchronous versions:

- `File.copyTo()` and `File.copyToAsync()`
- `File.deleteDirectory()` and `File.deleteDirectoryAsync()`
- `File.deleteFile()` and `File.deleteFileAsync()`
- `File.getDirectoryListing()` and `File.getDirectoryListingAsync()`
- `File.moveTo()` and `File.moveToAsync()`
- `File.moveToTrash()` and `File.moveToTrashAsync()`

Also, `FileStream` operations work synchronously or asynchronously depending on how the `FileStream` object opens the file: by calling the `open()` method or by calling the `openAsync()` method.

The asynchronous versions let you initiate processes that run in the background and dispatch events when complete (or when error events occur). Other code can execute while these asynchronous background processes are taking place. With asynchronous versions of the operations, you must set up event listener functions, using the `addEventListener()` method of the `File` or `FileStream` object that calls the function.

The synchronous versions let you write simpler code that does not rely on setting up event listeners. However, since other code cannot execute while a synchronous method is executing, important processes such as display object rendering and animation may be paused.

a. File Directories

a) Paths of File objects

Each File object has two properties that each defines its path:

Property	Description
nativePath	<p>Specifies the platform-specific path to a file.</p> <p>For example:</p> <p>Windows a path might be "c:\Sample directory\test.txt"</p> <p>On Mac OS it could be "/Sample directory/test.txt"</p> <p>A nativePath property uses the backslash (\) character as the directory separator character on Windows, and it uses the forward slash (/) character on Mac OS.</p>
url	<p>This may use the file URL scheme to point to a file.</p> <p>For example:</p> <p>Windows a path might be</p> <p>"file:///c:/Sample%20directory/test.txt"</p> <p>Mac OS it could be "file:///Sample%20directory/test.txt"</p>

b) Pointing a File object to a directory

There are different ways to set a File object to point to a directory.

c) Pointing to the user's home directory

You can point a File object to the user's home directory.

On Windows, the home directory is the parent of the "My Documents" directory

"C:\Documents and Settings\userName\My Documents").

On Mac OS, it is the Users/userName directory.

The following code sets a File object to point to an AIR Test subdirectory of the home directory:

```
var file:File = File.userDirectory.resolvePath("AIR Test");
```

d) Pointing to the user's documents directory

You can point a File object to the user's documents directory.

On Windows, the "My Documents" directory

"C:\Documents and Settings\userName\My Documents"

On Mac OS, it is the `Users/userName/Documents` directory.

The following code sets a File object to point to an AIR Test subdirectory of the documents directory:

e) Pointing to the desktop directory

You can point a File object to the desktop. The following code sets a File object to point to an AIR Test subdirectory of the desktop:

f) Pointing to the application storage directory

You can point a File object to the application storage directory. For every AIR application, there is a unique associated path that defines the application storage directory.

This directory is unique to each application and user. You may want to use this directory to store user-specific, application-specific data (such as user data or preferences files).

For example, the following code points a File object to a preferences file, `prefs.xml`, contained in the application storage directory:

```
var file:File = File.applicationStorageDirectory;  
file = file.resolvePath("prefs.xml");
```

The application storage directory location is based on the user name, the application ID, and the publisher ID:

On Mac OS—In:

```
/Users/user name/Library/Preferences/applicationID.publisherID/Local Store/
```

For example:

```
/Users/babbage/Library/Preferences/com.example.TestApp.02D88EEED35F8.1/Local Store
```

On Windows—In the documents and Settings directory, in:

```
username/Application Data/applicationID.publisherID/Local Store/
```

For example:

```
C:\Documents and
```

```
Settings\babbage\ApplicationData\com.example.TestApp.02D88EEED35F8.1\Local Store
```

g) Pointing to the application directory

You can point a File object to the directory in which the application was installed, known as the application directory. You can

reference this directory using the `File.applicationDirectory` property. You may use this directory to examine the application descriptor file or other resources installed with the application.

For example, the following code points a `File` object to a directory named `images` in the application directory:

```
var dir:File = File.applicationDirectory;
dir = dir.resolvePath("images");
```

The URL (and `url` property) for a `File` object created with `File.applicationDirectory` uses the app URL scheme:

```
var dir:File = File.applicationDirectory;
dir = dir.resolvePath("images");
trace(dir.url); // app:/images
```

h) Pointing to the filesystem root

The `File.getRootDirectories()` method lists all root volumes, such as `C:` and mounted volumes, on a Windows computer.

On Mac, this method always returns the unique root directory for the machine (the `/` directory).

i) Pointing to an explicit directory

You can point the `File` object to an explicit directory by setting the `nativePath` property of the `File` object, as in the following example (on Windows):

```
var file:File = new File();
file.nativePath = "C:\\AIR Test\\";
```

j) Navigating to relative paths

You can use the `resolvePath()` method to obtain a path relative to another given path.

For example, the following code sets a `File` object to point to an `"AIR Test"` subdirectory of the user's home directory:

```
var file:File = File.userDirectory;
file = file.resolvePath("AIR Test");
```

You can also use the `url` property of a `File` object to point it to a directory based on a URL string, as in the following:

```
var urlStr:String = "file:///C:/AIR Test/";
var file:File = new File()
file.url = urlStr;
```

k) Letting the user browse to select a directory

The `File` class includes the `browseForDirectory()` method, which presents a system dialog box in which the user can select a directory to assign to the object. The `browseForDirectory()` method is asynchronous.

It dispatches a `select` event if the user selects a directory and clicks the Open button, or it dispatches a `cancel` event if the user clicks the Cancel button.

For example, the following code lets the user select a directory and outputs the directory path upon selection:

```
var file:File = new File();
file.addEventListener(Event.SELECT, dirSelected);
file.browseForDirectory("Select a directory");

function dirSelected(e:Event):void {
    trace(file.nativePath);
}
```

l) Pointing to the directory from which the application was invoked

You can get the directory location from which an application is invoked, by checking the `currentDirectory` property of the `InvokeEvent` object dispatched when the application is invoked.

m) Creating directories

The `File.createDirectory()` method lets you create a directory. For example, the following code creates a directory named `AIR Test` as a subdirectory of the user's home directory:

```
var dir:File = File.userDirectory.resolvePath("AIR Test");
dir.createDirectory();
```

If the directory exists, the `createDirectory()` method does nothing. Also, in some modes, a `FileStream` object creates directories when opening files.

Missing directories are created when you instantiate a `FileStream` instance with the `fileMode` parameter of the `FileStream()` constructor set to `FileMode.APPEND` or `FileMode.WRITE`.

n) Creating a temporary directory

The `File` class includes a `createTempDirectory()` method, which creates a directory in the temporary directory folder for the System, as in the following example:

```
var temp:File = File.createTempDirectory();
```

The `createTempDirectory()` method automatically creates a unique temporary directory (saving you the work of determining a new unique location). You may use a temporary directory to temporarily store temporary files used for a session of the application.

Note that there is a `createTempFile()` method for creating new, unique temporary files in the System temporary directory. You may want to delete the temporary directory before closing the application, as it is not automatically deleted.

o) Enumerating directories

You can use the `getDirectoryListing()` method or the `getDirectoryListingAsync()` method of a `File` object to get an array of `File` objects pointing to files and subfolders in a directory.

For example, the following code lists the contents of the user's documents directory (without examining subdirectories):

```
var directory:File = File.documentsDirectory;
var contents:Array = directory.getDirectoryListing();

for (var i:uint = 0; i < contents.length; i++)
{
    trace(contents[i].name, contents[i].size);
}
```

p) Copying and moving directories

You can copy or move a directory, using the same methods as you would to copy or move a file.

For example, the following code copies a directory synchronously:

```
var sourceDir:File = File.documentsDirectory.resolvePath("AIR Test");
var resultDir:File = File.documentsDirectory.resolvePath("AIR Test Copy");
sourceDir.copyTo(resultDir);
```

When you specify `true` for the `overwrite` parameter of the `copyTo()` method, all files and folders in an existing target directory are deleted and replaced with the files and folders in the source directory (even if the target file does not exist in the source directory).

The directory that you specify as the `newLocation` parameter of the `copyTo()` method specifies the path to the resulting directory; it does not specify the parent directory that will contain the resulting directory.

q) Deleting directory contents

The `File` class includes a `deleteDirectory()` method and a `deleteDirectoryAsync()` method. These methods delete directories, the first working synchronously, the second working.

Both methods include a `deleteDirectoryContents` parameter (which takes a `Boolean` value); when this parameter is set to `true` (the default value is `false`) the call to the method deletes non-empty directories; otherwise, only empty directories are deleted.

For example, the following code synchronously deletes the `AIR Test` subdirectory of the user's documents directory:

```
var directory:File = File.documentsDirectory.resolvePath("AIR Test");
directory.deleteDirectory(true);
```

r) Copying and moving files

The `File` class includes two methods for copying files or directories: `copyTo()` and `copyToAsync()`.

The `File` class includes two methods for moving files or directories: `moveTo()` and `moveToAsync()`.

The `copyTo()` and `moveTo()` methods work synchronously, and the `copyToAsync()` and `moveToAsync()` methods work asynchronously.

To copy or move a file, you set up two `File` objects.

One points to the file to copy or move, and it is the object that calls the copy or move method; the other points to the destination (result) path.

The following copies a `test.txt` file from the `AIR Test` subdirectory of the user's documents directory to a file named `copy.txt` in the same directory:

```
var original:File = File.documentsDirectory.resolvePath("AIR Test/test.txt");
var newFile:File = File.resolvePath("AIR Test/copy.txt");
original.copyTo(newFile, true);
```

In this example, the value of `overwrite` parameter of the `copyTo()` method (the second parameter) is set to `true`. By setting this to `true`, an existing target file is overwritten. This parameter is optional.

If you set it to `false` (the default value), the operation dispatches an `IOErrorEvent` event if the target file exists (and the file is not copied).

The “Async” versions of the `copy` and `move` methods work asynchronously. Use the `addEventListener()` method to monitor

completion of the task or error conditions, as in the following code:

```
var original = File.documentsDirectory;
original = original.resolvePath("AIR Test/test.txt");

var destination:File = File.documentsDirectory;
destination = destination.resolvePath("AIR Test 2/copy.txt");

original.addEventListener(Event.COMPLETE, fileMoveCompleteHandler);
original.addEventListener(IOErrorEvent.IO_ERROR, fileMoveIOErrorHandler);
original.moveToAsync(destination);

function fileMoveCompleteHandler(event:Event):void
{
    trace(event.target); // [object File]
}

function fileMoveIOErrorHandler(event:IOErrorEvent):void
{
    trace("I/O Error.");
}
```

The File class also includes the File.moveToTrash() and File.moveToTrashAsync() methods, which move a file or directory to the system trash.

s) Deleting a file

The File class includes a deleteFile() method and a deleteFileAsync() method.

These methods delete files, the first working synchronously, the second working asynchronously.

For example, the following code synchronously deletes the test.txt file in the user's documents directory:

```
var file:File = File.documentsDirectory.resolvePath("test.txt");
file.deleteFile();
```

The following code asynchronously deletes the test.txt file of the user's documents directory:

```
var file:File = File.documentsDirectory.resolvePath("test.txt");
file.addEventListener(Event.COMPLETE, completeHandler)
file.deleteFileAsync();

function completeHandler(event:Event):void
```

```
{  
    trace("Deleted.")  
}
```

Also included are the `moveToTrash()` and `moveToTrashAsync()` methods, which you can use to move a file or directory to the System trash.

t) Moving a file to the trash

The `File` class includes a `moveToTrash()` method and a `moveToTrashAsync()` method.

These methods send a file or directory to the System trash, the first working synchronously, the second working asynchronously.

For example, the following code synchronously moves the `test.txt` file in the user's documents directory to the System trash:

```
var file:File = File.documentsDirectory.resolvePath("test.txt");  
file.moveToTrash();
```

u) Creating a temporary file

The `File` class includes a `createTempFile()` method, which creates a file in the temporary directory folder for the System, as in the following example:

```
var temp:File = File.createTempFile();
```

The `createTempFile()` method automatically creates a unique temporary file (saving you the work of determining a new unique location).

You may use a temporary file to temporarily store information used in a session of the application. Note that there is also a `createTempDirectory()` method, for creating a unique temporary directory in the System temporary directory.

You may want to delete the temporary file before closing the application, as it is not automatically deleted.

v) FileStream open modes

The `open()` and `openAsync()` methods of a `FileStream` object each include a `fileMode` parameter, which defines some properties for a file stream, including the following:

- The ability to read from the file
- The ability to write to the file

- Whether data will always be appended past the end of the file (when writing)
- What to do when the file does not exist (and when its parent directories do not exist)

The following are the various file modes (which you can specify as the `fileMode` parameter of the `open()` and `openAsync()` methods):

File Mode	Description
<code>FileMode.READ</code>	Specifies that the file is open for reading only.
<code>FileMode.WRITE</code>	Specifies that the file is open for writing. If the file does not exist, it is created when the <code>FileStream</code> object is opened. If the file does exist, any existing data is deleted.
<code>FileMode.APPEND</code>	
<code>FileMode.UPDATE</code>	

5.3 - Create and customize native windows and menus.

a. Windows in AIR

AIR supports three distinct APIs for working with windows: the ActionScript-oriented `NativeWindow` class, the Flex framework `mx:WindowedApplication` and `mx:Window` classes, which "wrap" the `NativeWindow` class, and, in the HTML environment, the JavaScript `Window` class.

a) ActionScript windows

When you create windows with the `NativeWindow` class, you use the Flash stage and display list directly. To add a visual object to a `NativeWindow`, you add the object to the display list of the window stage or to another display object on the stage.

b) Flex Framework windows

When you create windows with the Flex framework, you typically use MXML components to populate the window. To add a Flex component to a window, you add the component element to the window MXML definition.

You can also use ActionScript to add content dynamically. The `mx:WindowedApplication` and `mx:Window` components are designed as Flex containers and so can accept Flex components directly whereas `NativeWindow` objects cannot.

When necessary, `NativeWindow` properties and methods can be accessed through the `WindowedApplication` and `Window` objects using the `stage.nativeWindow` property. See [About window containers](#) for more information about these Flex components.

c) HTML windows

When you create HTML windows, you use HTML, CSS, and JavaScript to display content. To add a visual object to an HTML window, you add that content to the HTML DOM. HTML windows are a special category of `NativeWindow`.

The AIR host defines a `nativeWindow` property in HTML windows that provides access to the underlying `NativeWindow` instance. You can use this property to access the `NativeWindow` properties, methods, and events described in this section.

Note: The JavaScript `Window` object also has methods for scripting the containing window, such as `moveTo()` and `close()`. Where overlapping methods are available, you can use the method most convenient.

d) The initial application window

The first window of your application is automatically created for you by AIR. AIR sets the properties and content of the window using the parameters specified in the `initialWindow` element of the application descriptor file.

If the root content is a SWF file, AIR creates a `NativeWindow` instance, loads the SWF file, and adds it to the window stage. If the root content is an HTML file, AIR creates an HTML window and loads the HTML.

b. Properties controlling native window style and behavior

The following properties control the basic appearance and behavior of a window:

- `type`
- `systemChrome`
- `transparent`

When you create a window, you set these properties on the `NativeWindowInitOptions` object passed to the window constructor. AIR reads the properties for the initial application window from the application descriptor. (Except the `type` property, which cannot be set in the application descriptor and is always set to `normal`.) The properties cannot be changed after window creation.

Some settings of these properties are mutually incompatible:

`systemChrome` cannot be set to `standard` when either `transparent` is `true` or `type` is `lightweight`.

a) Window Types

The AIR window types combine chrome and visibility attributes of the native operating system to create three functional types of window. Use the constants defined in the `NativeWindowType` class to reference the type names in code. AIR provides the following window types:

Type	Description
normal	A typical window. Normal windows use the full-size style of chrome and appear on the Windows task bar and the Mac OS X window menu.
utility	A tool palette. Utility windows use a slimmer version of the system chrome and do not appear on the Windows task bar and the Mac OS-X window menu.
lightweight	<p>Lightweight windows have no chrome and do not appear on the Windows task bar or the Mac OS X window menu.</p> <p>In addition, lightweight windows do not have the System (Alt+Space) menu on Windows.</p> <p>Lightweight windows are suitable for notification bubbles and controls such as combo-boxes that open a short-lived display area.</p> <p>When the lightweight type is used, <code>systemChrome</code> must be set to <code>none</code>.</p>

b) Window chrome

Window chrome is the set of controls that allow users to manipulate a window in the desktop environment.

Chrome elements include the title bar, title bar buttons, border, and resize grippers.

c) System chrome

You can set the `systemChrome` property to `standard` or `none`.

Choose `standard` system chrome to give your window the set of standard controls created and styled by the user's operating system.

Choose `none` to provide your own chrome for the window (or to use Flex chrome). Use the constants defined in the `NativeWindowSystemChrome` class to reference the system chrome settings in code.

System chrome is managed by the system. Your application has no direct access to the controls themselves, but can react to the

events dispatched when the controls are used. When you use standard chrome for a window, the `transparent` property must be set to `false` and the `type` property must be `normal` or `utility`.

d) Flex chrome

When you use the Flex `mx:WindowedApplication` or `mx:Window` components, the window can be use either system chrome or chrome provided by the Flex framework. To use the Flex chrome, set the `systemChrome` property used to create the window to `none`.

e) Custom chrome

When you create a window with no system chrome and you do not use the Flex `mx:WindowedApplication` or `mx:Window` components, then you must add your own chrome controls to handle the interactions between a user and the window.

You are also free to make transparent, non-rectangular windows.

f) Window transparency

To allow alpha blending of a window with the desktop or other windows, set the window `transparent` property to `true`.

The `transparent` property must be set before the window is created and cannot be changed.

A transparent window has no default background. Any window area not occupied by a display object will be invisible.

If a display object has an alpha setting of less than one, then anything below the object will show through, including other display objects in the same window, other windows, and the desktop.

Rendering large alpha-blended areas can be slow, so the effect should be used conservatively.

Transparent windows are useful when you want to create applications with borders that are irregular in shape or that "fade out" or appear to be invisible.

Transparency cannot be used with windows that have system chrome.

c. Creating windows

AIR automatically creates the first window for an application, but you can create any additional windows you need.

To create a native window, use the `NativeWindow` constructor method.

To create a Flex window, use the `mx:Window` class.

To create an HTML window, either use the `HTMLLoader.createRootWindow()` method or, from an HTML document, call the `JavaScriptWindow.open()` method.

a) Specifying window initialization properties

Set the properties for the initial window created by AIR in the application descriptor file.

The main window of an AIR application is always type, `normal`.

(Additional window properties can be specified in the descriptor file, such as `visible`, `width`, and `height`, but these properties can be changed at any time.)

When you create a window with the Flex `mx:Window` class, specify the initialization properties on the window object itself, either in the MXML declaration for the window, or in the code that creates the window.

The desktop window is not created until you call the window `open()` method. Once a window is opened, these initialization properties cannot be changed.

Set the properties for other native and HTML windows created by your application using the `NativeWindowInitOptions` class.

When you create a window, you must pass a `NativeWindowInitOptions` object specifying the window properties to either the `NativeWindow` constructor function or the `HTMLLoader.createRootWindow()` method.

The following code creates a `NativeWindowInitOptions` object for a utility window:

```
var options:NativeWindowInitOptions = new NativeWindowInitOptions();
options.systemChrome = NativeWindowSystemChrome.STANDARD;
options.type = NativeWindowType.UTILITY
options.transparent = false;
options.resizable = false;
options.maximizable = false;
```

Setting `systemChrome` to `standard` when `transparent` is `true` or `type` is `lightweight` is not supported.

b) Creating the initial application window

AIR creates the initial application window based on the properties specified in the application descriptor and loads the file referenced in the content element. The content must be a SWF or an HTML file.

The initial window can be the main window of your application or it can merely serve to launch one or more other windows. You do not have to make it visible at all.

c) Creating the initial window with Flex

When creating an AIR application with the Flex framework, use the `mx:WindowedApplication` as the root element of your main MXML file.

(You can use the `mx:Application` component, but it does not support all the window features available in AIR.)

The `WindowedApplication` component serves as the initial entry point for the application.

When you launch the application, AIR creates a native window, initializes the Flex framework, and adds the `WindowedApplication` object to the window stage.

When the launch sequence finishes, the `WindowedApplication` dispatches an `applicationComplete` event. Access the desktop window object with the `nativeWindow` property of the `WindowedApplication` instance.

d) Creating the initial window with ActionScript

When you create an AIR application using the Flex 3 SDK and ActionScript, the main class of your application must extend the `Sprite` class (or a subclass of the `Sprite` class).

This class serves as the main entry point for the application.

When your application launches, AIR creates a window, creates an instance of the main class, and adds the instance to the window stage.

To access the window, you can listen for the `addedToStage` event and then use the `nativeWindow` property of the `Stage` object to get a reference to the `NativeWindow` object.

e) Creating a new NativeWindow

To create a new `NativeWindow`, pass a `NativeWindowInitOptions` object to the `NativeWindow` constructor:

```
var options:NativeWindowInitOptions = new NativeWindowInitOptions();
options.systemChrome = NativeWindowSystemChrome.STANDARD;
options.transparent = false;

var newWindow:NativeWindow = new NativeWindow(options);
```

The window is not shown until you set the `visible` property to `true` or call the `activate()` method

Once the window is created, you can initialize its properties and load content into the window using the stage property and Flash display list techniques.

In almost all cases, you should set the stage `scaleMode` property of a new native window to `noScale` (use the `StageScaleMode.NO_SCALE` constant).

The Flash scale modes are designed for situations in which the application author does not know the aspect ratio of the

application display space in advance.

The scale modes let the author choose the least-bad compromise: clip the content, stretch or squash it, or pad it with empty space.

Since you control the display space in AIR (the window frame), you can size the window to the content or the content to the window without compromise.

The scale mode for Flex and HTML windows is set to `noScale` automatically.

```
var maxOSSize:Point = NativeWindow.systemMaxSize;  
var minOSSize:Point = NativeWindow.systemMinSize;
```

f) Creating a new HTML window

To create a new HTML window, you can either call the JavaScript `window.open()` function, or you can call the AIR `HTMLLoader` class `createRootWindow()` function.

HTML content in any security sandbox can use the standard JavaScript `Window.open()` method.

If the content is running outside the application sandbox, the `open()` method can only be called in response to user interaction, such as a mouse click or key press.

When `open()` is called, a window with system chrome will be created to display the content at the specified URL.

For example:

```
newWindow = window.open("xml.html", "logWindow", "height=600, width=400, top=10, left=10");
```

g) Adding content to a window

How you add content to an AIR window depends on the type of window.

MXML and HTML let you declaratively define the basic content of the window. You can embed resources in the application SWF or you can load them from separate application files. Flex, Flash and HTML content can all be created on the fly and added to a window dynamically.

When you load SWF content, or HTML content containing JavaScript, you must take the AIR security model into consideration.

Any content in the application security sandbox, that is, content installed with your application and loadable with the `app:` URL scheme, will have full privileges to access all the AIR APIs.

Any content loaded from outside this sandbox cannot access the AIR APIs. JavaScript content outside the application sandbox will not be able to use the runtime, `nativeWindow` or `htmlLoader` properties of the JavaScript `Window` object.

To allow safe cross-scripting, you can use a sandbox bridge to provide a limited interface between application content and non-application content.

In HTML content, you can also map pages of your application into a non-application sandbox to allow the code on that page to cross-script external content. See AIR security.

d. Loading a SWF or image

You can load Flash or images into the display list of a native window using the `flash.display.Loader` class.

```
package {
    import flash.display.Sprite;
    import flash.events.Event;
    import flash.net.URLRequest;
    import flash.display.Loader;

    public class LoadedSWF extends Sprite
    {
        public function LoadedSWF() {
            var loader:Loader = new Loader();
            loader.load(new URLRequest("visual.swf"));
            loader.contentLoaderInfo.addEventListener(Event.COMPLETE, loadFlash);
        }

        private function loadFlash(event:Event):void {
            addChild(event.target.loader);
        }
    }
}
```

You can load a SWF file that contains library code for use in an HTML-based application. The simplest way to load a SWF in an HTML window is to use the script tag, but you can also use the Loader API directly.

a) Adding Flash content as an overlay on an HTML window

Because HTML windows are contained within a `NativeWindow` instance, you can add Flash display objects both above and below the HTML layer in the display list.

To add a display object above the HTML layer, use the `addChild()` method of the `window.nativeWindow.stage` property. The `addChild()` method will add content layered above any existing content in the window.

To add a display object below the HTML layer, use the `addChildAt()` method of the `window.nativeWindow.stage` property, passing in a value of zero for the `index` parameter.

Placing an object at the zero index will move existing content, including the HTML display, up one layer and insert the new content at the bottom.

For content layered underneath the HTML page to be visible, you must set the `paintsDefaultBackground` property of the `HTMLLoader` object to `false`. In addition, any elements of the page that set a background color will not be transparent.

If, for example, you set a background color for the body element of the page, none of the page will be transparent.

b) Loading HTML content into a NativeWindow

To load HTML content into a `NativeWindow`, you can either add an `HTMLLoader` object to the window stage and load the HTML content into the `HTMLLoader`, or create a new window that already contains an `HTMLLoader` object by using the `HTMLLoader.createRootWindow()` method.

The following example displays HTML content within a 300 by 500 pixel display area on the stage of a native window:

//newWindow is a NativeWindow instance

```
var htmlView:HTMLLoader = new HTMLLoader();
html.width = 300;
html.height = 500;
```

//set the stage so display objects are added to the top-left and not scaled

```
newWindow.stage.align = "TL";
newWindow.stage.scaleMode = "noScale";
newWindow.stage.addChild( htmlView );
```

//urlString is the URL of the HTML page to load

```
htmlView.load( new URLRequest(urlString) );
```

b. AIR supports the following types of menus:

a) Application Menu

An application menu is a global menu that applies to the entire application. Application menus are supported on Mac OS X, but not on Windows.

On Mac OS X, the operating system automatically creates an application menu. You can use the AIR menu API to add items and submenus to the standard menus. You can add listeners for handling the existing menu commands. Or you can remove existing items.

b) Window menus

A window menu is associated with a single window and is displayed below the title bar. Menus can be added to a window by creating a new `NativeMenu` object and assigning it to the `menu` property of the `NativeWindow` object. Window menus are supported on the Windows operating system, but not on Mac OS X. Native menus can only be used with windows that have system chrome.

c) Context menus

Context menus open in response to a right-click or command-click on an interactive object in SWF content or a document element in HTML content.

You can create a context menu using the `AIRNativeMenu` class. (You can also use the legacy Flash `ContextMenu` class.) In HTML content, you can use the Webkit HTML and JavaScript APIs to add context menus to an HTML element.

d) Dock and system tray icon menus

These icon menus are similar to context menus and are assigned to an application icon in the Mac OS X dock or Windows notification area. Dock and system tray icon menus use the `NativeMenu` class.

On Mac OS X, the items in the menu are added above the standard operating system items. On Windows, there is no standard menu.

e) Pop-up menus

An AIR pop-up menu is like a context menu, but is not necessarily associated with a particular application object or component. Pop-up menus can be displayed anywhere in a window by calling the `display()` method of any `NativeMenu` object.

f) Flex menus

The Flex framework provides a set of Flex menu components. The Flex menus are drawn by the AIR runtime rather than the operating system and are not native menus.

A Flex menu component can be used for Flex windows that do not have system chrome.

Another benefit of using the Flex menu component is that you can specify menus declaratively in MXML format. If you are using the Flex Framework, use the Flex menu classes for window menus instead of the native classes.

g) Custom menus

Native menus are drawn entirely by the operating system and, as such, exist outside the Flash and HTML rendering models.

You are of course free to create your own non-native menus using MXML, ActionScript, or JavaScript. The AIR menu classes do not provide any facility for controlling the drawing of native menus.

h) Default menus

The following default menus are provided by the operating system or the built-in AIR class:

- Application menu on Mac OS X
- Dock icon menu on Mac OS X
- Context menu for selected text and images in HTML content
- Context menu for selected text in a `TextField` object (or an object that extends `TextField`)

i) Menu structure

Menus are hierarchical in nature. `NativeMenu` objects contain child `NativeMenuItem` objects. `NativeMenuItem` objects that represent submenus, in turn, can contain `NativeMenu` objects.

The top- or root-level menu object in the structure represents the menu bar for application and window menus. (Context, icon, and pop-up menus don't have a menu bar).

j) Menu events

`NativeMenu` and `NativeMenuItem` objects both dispatch displaying and select events:

- `displaying`: Immediately before a menu is displayed, the menu and its menu items will dispatch a displaying event to any registered listeners.

The displaying event gives you an opportunity to update the menu contents or item appearance before it is shown to the user.

For example, in the listener for the displaying event of an "Open Recent" menu, you could change the menu items to reflect the current list of recently viewed documents.

The `target` property of the event object is always the menu that is about to be displayed. The `currentTarget` is the object on which the listener is registered, either the menu itself, or one of its items.

Note: The displaying event is also dispatched whenever the state of the menu or one of its items is accessed.

- `select`: When a command item is chosen by the user, the item will dispatch a select event to any registered listeners. Submenu and separator items cannot be selected and so never dispatch a select event.

A `select` event bubbles up from a menu item to its containing menu, on up to the root menu. You can listen for `select` events directly on an item and you can listen higher up in the menu structure.

When you listen for the `select` event on a menu, you can identify the selected item using the event `target` property.

As the event bubbles up through the menu hierarchy, the `currentTarget` property of the event object identifies the current menu object.

Note: `ContextMenu` and `ContextMenu` objects dispatch `menuItemSelect` and `menuSelect` events as well as `select` and displaying events.

k) Menu item state

Menu items have the two state properties, `checked` and `enabled`:

- `checked`: Set to `true` to display a check mark next to the item label.
- `enabled`: Toggle the value between `true` and `false` to control whether the command is enabled. Disabled items are visually "grayed-out" and do not dispatch `select` events.

b. Creating Menus

a) Creating a root menu object

To create a new `NativeMenu` object to serve as the root of the menu, use the `NativeMenu` constructor:

```
var root:NativeMenu = new NativeMenu();
```

For application and window menus, the root menu represents the menu bar and should only contain items that open submenus.

Context menu and pop-up menus do not have a menu bar, so the root menu can contain commands and separator lines as well as submenus.

After the menu is created, you can add menu items. Items appear in the menu in the order in which they are added, unless you add the items at a specific index using the `addItemAt()` method of a menu object.

To assign the menu as an application, window, icon, or context menu, or display it as a pop-up menu as shown in the following sections:

b) Setting the application menu

```
NativeApplication.nativeApplication.menu = root;
```

Note: Mac OS X defines a menu containing standard items for every application. Assigning a new `NativeMenu` object to the `menu` property of the `NativeApplication` object will replace the standard menu. You can also use the standard menu instead of replacing it.

c) Setting a window menu

```
nativeWindowObject.menu = root;
```

d) Setting a context menu on an interactive object

```
interactiveObject.contextMenu = root;
```

e) Setting a dock icon menu

```
DockIcon(NativeApplication.nativeApplication.icon).menu = root;
```

Note: Mac OS X defines a standard menu for the application dock icon. When you assign a new `NativeMenu` to the `menu` property of the `DockIcon` object, the items in that menu are displayed above the standard items. You cannot remove, access, or

modify the standard menu items.

f) Setting a system tray icon menu

```
SystemTrayIcon(NativeApplication.nativeApplication.icon).menu = root;
```

g) Displaying a menu as a pop-up

```
root.display(stage, x, y);
```

h) Creating a submenu

To create a submenu, you add a `NativeMenuItem` object to the parent menu and then assign the `NativeMenu` object defining the submenu to the item's `submenu` property. AIR provides two ways to create submenu items and their associated menu object:

You can create a new menu item and its related menu object in one step with the `addSubmenu()` method:

```
var editMenuItem:NativeMenuItem = root.addSubmenu(new NativeMenu(), "Edit");
```

You can also create the menu item and assign the menu object to its `submenu` property separately:

```
var editMenuItem:NativeMenuItem = root.addItem("Edit", false);  
editMenuItem.submenu = new NativeMenu();
```

i) Creating a menu command

To create a menu command, add a `NativeMenuItem` object to a menu and add an event listener referencing the function implementing the menu command:

```
var copy:NativeMenuItem = new NativeMenuItem("Copy", false);  
copy.addEventListener(Event.SELECT, onCopyCommand);  
editMenu.addItem(copy);
```

You can listen for the select event on the command item itself (as shown in the example), or you can listen for the select event on a parent menu object.

Note: Menu items that represent submenus and separator lines do not dispatch select events and so cannot be used as commands.

j) Creating a menu separator line

To create a separator line, create a `NativeMenuItem`, setting the `isSeparator` parameter to `true` in the constructor. Then add the separator item to the menu in the correct location:

```
var separatorA:NativeMenuItem = new NativeMenuItem("A", true);
    editMenu.addItem(separatorA);
```

Note: that the label specified for the separator, if any, will not be displayed.

5.4 - Adding drag-and-drop functionality to and from the desktop.

a. Drag-and-drop gesture stages

The drag-and-drop gesture has three stages:

- **Initiation:** A user initiates a drag-and-drop operation by dragging from a component, or an item in a component, while holding down the mouse button.

The component that is the source of the dragged item is typically designated as the **drag initiator and dispatches `nativeDragStart` and `nativeDragComplete` events.**

An Adobe AIR application starts a drag operation by calling the **`NativeDragManager.doDrag()`** method in response to a `mouseDown` or `mouseMove` event.

- **Dragging:** While holding down the mouse button, the user moves the mouse cursor to another component, application, or to the desktop. AIR optionally displays a proxy image during the drag. As long as the drag is underway, the initiator **object dispatches `nativeDragUpdate` events.**

When the user moves the mouse over a possible drop target in an AIR application, **the drop target dispatches a `nativeDragEnter` event.**

The event handler can inspect the event object to determine whether the dragged data is available in a format that the target accepts and, if so, let the user drop the data onto it by calling the **`NativeDragManager.acceptDragDrop()`** method.

As long as the drag gesture remains over an interactive object, that **object dispatches `nativeDragOver` events.** When the drag gesture leaves the interactive object, it **dispatches a `nativeDragExit` event.**

- **Drop:** The user releases the mouse over an eligible drop target. If the target is an AIR application or component, then the **component dispatches a `nativeDragDrop` event.**

The event handler can access the transferred data from the event object. If the target is outside AIR, the operating system or another application handles the drop.

In both cases, the initiating object dispatches a `nativeDragComplete` event (if the drag started from within AIR). The `NativeDragManager` class controls both drag-in and drag-out gestures. All the members of the `NativeDragManager` class are static; do not create an instance of this class.

a) The Clipboard object

Data that is dragged into or out of an application or component is contained in a `Clipboard` object. A single `Clipboard` object can make available different representations of the same information to increase the likelihood that another application can understand and use the data.

For example, an image could be included as image data, a serialized `Bitmap` object, and as a file. Rendering of the data in a format can be deferred to a rendering function that is not called until the data is read.

Once a drag gesture has started, the `Clipboard` object can only be accessed from within an event handler for the `nativeDragEnter`, `nativeDragOver`, and `nativeDragDrop` events.

After the drag gesture has ended, the `Clipboard` object cannot be read or reused. An application object can be transferred as a reference and as a serialized object.

References are only valid within the originating application. Serialized object transfers are valid between AIR applications, but can only be used with objects that remain valid when serialized and deserialized.

Objects that are serialized are converted into the `ActionMessageFormat` for `ActionScript 3 (AMF3)`, a string-based data-transfer format.

b) Working with the Flex framework

In most cases, it is better to use the Adobe Flex drag-and-drop API when building Flex applications. The Flex framework provides an equivalent feature set when a Flex application is run in AIR (it uses the `AIRNativeDragManager` internally).

Flex also maintains a more limited feature set when an application or component is running within the more restrictive browser environment. AIR classes cannot be used in components or applications that run outside the AIR run-time environment.

c) Starting a drag-out operation

To start a drag operation, call the `NativeDragManager.doDrag()` method in response to a mouse down event.

The `doDrag()` method is a static method that takes the following parameters:

d) Completing a drag-out transfer

When a user drops the dragged item by releasing the mouse, the initiator object dispatches a `nativeDragComplete` event. You can check the `dropAction` property of the event object and then take the appropriate action.

For example, if the action is `NativeDragAction.MOVE`, you could remove the source item from its original location. The user can abandon a drag gesture by releasing the mouse button while the cursor is outside an eligible drop target.

The drag manager sets the `dropAction` property or an abandoned gesture to `NativeDragAction.NONE`.

e) Supporting the drag-in gesture

To support the drag-in gesture, your application (or, more typically, a visual component of your application) must respond to `nativeDragEnter` or `nativeDragOver` events. Steps in a typical drop operation

The following sequence of events is typical for a drop operation:

- The user drags a clipboard object over a component.
- The component dispatches a `nativeDragEnter` event.
- The `nativeDragEnter` event handler examines the event object to check the available data formats and allowed actions. If the component can handle the drop, it calls `NativeDragManager.acceptDragDrop()`.
- The `NativeDragManager` changes the mouse cursor to indicate that the object can be dropped.
- The user drops the object over the component.
- The receiving component dispatches a `nativeDragDrop` event.
- The receiving component reads the data in the desired format from the `Clipboard` object within the event object.
- If the drag gesture originated within an AIR application, then the initiating interactive object dispatches a `nativeDragComplete` event. If the gesture originated outside AIR, no feedback is sent.

5.5 - Install, uninstall, and update an AIR application.

a. Installing, Updating and Uninstalling Air Applications

a) Installing

AIR applications are distributed via AIR installer files, which use the air extension. When Adobe AIR is installed and an AIR file is opened, the runtime administers and manages the application installation process.

Note: Developers can specify a version, and application name, and a publisher source, but the initial application installation workflow itself cannot be modified.

This restriction is advantageous for users because all AIR applications share a secure, streamlined, and consistent installation procedure administered by Adobe AIR. If application customization is necessary, it can be provided when the application is first executed.

The default application installer provides the user with security-related information. **AIR displays the publisher name during installation when the AIR application has been signed with a certificate that is trusted, or which chains to a certificate that is trusted on the installation computer.** Otherwise the publisher name is displayed as "Unknown." This lets the user make an informed decision whether to install the application or not:

AIR applications first require the runtime to be installed on a user's computer, just as SWF files first require the Flash Player browser

plug-in to be installed.

The runtime can be installed in two ways:

- Using the seamless install feature
- Or via a manual installation.

The seamless install feature provides developers with a streamlined installation experience for users who do not have Adobe AIR installed yet. In the seamless install method, the developer embeds a SWF file in a web page, and that SWF file presents the name of the AIR application for installation. When a user clicks in the SWF file to install the application, the SWF file checks for the presence of the runtime. If the runtime cannot be detected it is installed, and the runtime is activated immediately with the installation process for the developer's application. The user is provided with the option to cancel installation.

Alternatively, the user can manually download and install the runtime before installing an AIR file. The developer can then distribute an AIR file by different means (for example, via e-mail or an HTML link on a website). When the AIR file is opened, the runtime is activated and begins to process the application installation.

The AIR security model allows users to decide whether to install an AIR application. The AIR installer provides several improvements over native application install technologies that make this trust decision easier for users:

The runtime provides a consistent installation experience on all operating systems, even when an AIR application is installed from a link in a web browser. Most native application install experiences depend upon the browser or other application to provide security information, if it is provided at all.

The AIR application installer identifies the source of the application (or, if the source cannot be verified, the installer makes this clear) and it provides information about the privileges that are available to the application if the user allows the installation to proceed.

The runtime administers the installation process of an AIR application. An AIR application cannot manipulate the installation process the runtime uses.

In general, users should not install any desktop application (including an AIR application) that comes from a source that they do not trust, or that cannot be verified. The burden of proof on security for native applications is equally true for AIR applications as it is for other installable applications

b) Updating AIR applications

Development and deployment of software updates are one of the biggest security challenges facing native code applications.

An installed AIR application can check a remote location for an update AIR file. If an update is appropriate, the AIR file is downloaded, installed, and the application restarts.

The developer documentation provides details on using this method not only to provide new functionality but also respond to potential security vulnerabilities.

c) Removing an AIR application

A user can remove an AIR application:

- On Windows: Using the Add/Remove Programs panel to remove the application.
- On Mac OS: Deleting the application file from the install location.

Removing an AIR application removes all files in the application directory. However, it does not remove files that the application may have written outside the application directory. Removing AIR applications does not revert changes the AIR application has made to files outside the application directory.

5.6 - List and describe the AIR security contexts.

a. Security Contexts

a) Code Signing

Adobe AIR requires all AIR applications to be digitally signed. Code signing is a process of digitally signing code to ensure integrity of software and the identity of the publisher. Developers can sign AIR applications with a certificate issued by a Certification Authority (CA) or by constructing a self-signed certificate.

Digitally signing AIR files with a certificate issued by a recognized certificate authority (CA) provides significant assurance to users that the application they are installing has not been accidentally or maliciously altered.

Digitally signing AIR files with a certificate issued by a recognized certificate authority (CA) identifies the developer as the signer (publisher). AIR recognizes code-signing certificates issued by the VeriSign and Thawte certificate authorities.

The AIR application installer displays the publisher name during installation when the developer has signed the AIR file with a VeriSign or Thawte certificate.

The AIR application installer displays the publisher name during installation when the AIR application has been signed with a certificate that is trusted, or which chains to a certificate that is trusted on the installation computer. The Certification Authority (CA) verifies the publisher or developer's identity using established verification processes before issuing a high assurance certificate.

Developers can also sign AIR applications using a self-signed certificate; one that they create themselves. However, the AIR application installer presents these applications as originating from an unverified publisher.

When an AIR file is signed, a digital signature is included in the installation file. The signature includes a digest of the package, which is used to verify that the AIR file has not been altered since it was signed, and it includes information about the signing certificate, which is used to verify the publisher identity.

AIR uses the public key infrastructure (PKI) supported through the operating system's certificate store. The computer on which an AIR application is installed must either directly trust the certificate used to sign the AIR application, or it must trust a chain of certificates linking the certificate to a trusted certificate authority in order for the publisher information to be verified.

If an AIR file is signed with a certificate that does not chain to one of the trusted root certificates (and normally this includes all

self-signed certificates), then the publisher information cannot be verified. While AIR can determine that the AIR file has not been altered since it was signed, there is no way to verify who actually created and signed the file.

b) Security sandboxes

AIR provides a comprehensive security architecture that defines permissions for each file in an AIR application. This includes both those files installed with the application and other files loaded by the application.

Permissions are granted to files according to their origin, and are assigned to logical security groupings called sandboxes.

Files installed with the application are in a directory known as the application directory, and as such, they are, by default, placed in a security sandbox — known as the application sandbox — that has access to all AIR APIs.

This includes APIs that would pose a great security risk if made available to content from sources other than the application resource directory (in other words, files that are not installed with the application).

The AIR security model of sandboxes is composed of the Flash Player security model with the addition of the application sandbox. Files that are not in the application sandbox have security restrictions like those specified by the Flash Player security model.

The runtime uses these security sandboxes to define the range of data that a file may access and the operations it may execute. To maintain local security, the files in each sandbox are isolated from the files of other sandboxes.

For example, a SWF file loaded into an AIR application from an external Internet URL is placed into the remote sandbox, and does not by default have permission to script into files that reside in the application directory, which are assigned to the application sandbox

c) Accessing the file system

Applications running in a web browser have only limited interaction with the user's local file system. Web browsers implement security policies that ensure that a user's computer cannot be compromised as a result of loading web content.

For example, SWF files running through Flash Player in a browser cannot directly interact with files already on a user's computer.

Shared objects can be written to a user's computer for the purpose of maintaining user preferences and other data, but this is the limit of file system interaction.

Because AIR applications are natively installed, they have a different security contract with the end user. This contract between the application and the end user is made at install time just like native applications, and it includes the capability for the application to read and write across the local file system.

This freedom comes with a higher degree of responsibility for developers. Accidental application security gaps jeopardize not only the functionality of the application, but also the integrity of the user's computer. The developer documentation includes an "AIR Security" chapter that addresses best practices.

Unless there are administrator restrictions applied to the user's computer, AIR applications are privileged to write to any location on the user's hard drive. However, developers are encouraged to use the user- and application-specific application storage

directory that the runtime provides for each application.

The AIR API provides convenient methods for developers to read and write data in the application storage directory. The runtime also provides an encrypted local data storage area unique to each application and user.

This allows applications to save and retrieve data that is stored on the user's local hard drive in an encrypted format that cannot be deciphered by other applications or users. A separate encrypted local store is used for each AIR application, and each AIR application uses a separate encrypted local store for each user.

Applications may use the encrypted local store to store information that must be secured, such as login credentials for web services. AIR uses DPAPI on Windows and KeyChain on Mac OS to associate encrypted local stores to each user. The encrypted local store uses AES-CBC 128-bit encryption.

d) Working securely with un-trusted content

Content not assigned to the application sandbox can provide additional scripting functionality to an AIR application, but only if it meets the security criteria of the runtime. This section explains the AIR security contract with non-application content.

AIR applications restrict scripting access for non-application content more stringently than the Flash Player browser plug-in restricts scripting access for un-trusted content. For example, in Flash Player in the browser, a SWF file can call the `System.allowDomain()` method to grant scripting access to any SWF content loaded from a specified domain.

Calls to this method are not permitted for content in the AIR application security sandbox, since it would grant unreasonable access to the non-application file into the user's file system.

AIR applications that script between application and non-application content have more complex security arrangements. Files that are not in the application sandbox are only allowed to access the properties and methods of files in the application sandbox through the use of a sandbox bridge.

A sandbox bridge acts as a gateway between application content and non-application content, providing explicit interaction between the two files. When used correctly, sandbox bridges provide an extra layer of security, restricting non-application content from accessing object references that are part of application content.

The benefit of sandbox bridges is best illustrated through example. Suppose an AIR music store application wants to provide an API to advertisers who want to create their own SWF files, with which the store application can then communicate. The store wants to provide advertisers with methods to look up artists and CDs from the store, but also wants to isolate some methods and properties from the third-party SWF file for security reasons.

A sandbox bridge can provide this functionality. By default, content loaded externally into an AIR application at runtime does not have access to any methods or properties in the main application.

With a custom sandbox bridge implementation, a developer can provide services to the remote content without exposing these methods or properties. The sandbox bridge provides a limited pathway between trusted and un-trusted content.

The "AIR Security" chapter of the developer documentation provide full details on using sandbox bridges.

e) HTML security

HTML content has different security considerations than SWF-based content, primarily due to the ability of JavaScript to create dynamically generated code.

Dynamically generated code, such as that which is made when calling the `eval()` function, could pose a security risk if allowed within the application sandbox. For example, an application could inadvertently execute a string loaded from a network sandbox, and that string may contain malicious code, such as code to delete or alter files on the user's computer or to report back the contents of a local file to an un-trusted network domain.

Ways to generate dynamic code include the following:

- Calling the `eval()` function.
- Setting `innerHTML` properties or calling DOM functions to insert script tags to load a script outside the resource directly.
- Setting `innerHTML` properties or calling DOM functions to insert script tags that have inline code (rather than loading a script via the `src`).
- Setting the `src` for script tags for content in the application sandbox to a file that is not in the application resource directory.
- Using the javascript URL scheme (as in `href="javascript:alert('Test')"`).

Code in the application security sandbox can only use these methods while content is loading from application directory. This prevents code in the application sandbox, which has access to the full AIR APIs, from executing scripts from potentially un-trusted sources.

b. Other security considerations

Although AIR applications are built using web technologies, it is important for developers to note that they are not working within the browser security model. This means that it is possible to build AIR applications that can do harm to the local system, either intentionally or unintentionally. AIR attempts to minimize this risk, but there are still ways where vulnerabilities can be introduced. This section covers important potential insecurities. The developer documentation provides best practices for building applications that avoid these risks. Risk from importing files into the application security sandbox

Content the application sandbox has the full privileges of the runtime. Developers are advised to consider the following:

- Include a file in an AIR file (in the installed application) only if it is necessary.
- Include a scripting file in an AIR file (in the installed application) only if its behavior is fully understood and trusted.

Do not use data from a network source as parameters to methods of the AIR API that may lead to code execution. Adobe AIR protects content in the application sandbox from using data from a network source as code, which could lead to inadvertent execution of malicious code. This includes use of the `ActionScript Loader.loadBytes()` method and the JavaScript `eval()`

function.

a) Risk from using an external source to determine paths

An AIR application can be compromised when using external data or content. For this reason, applications must take special care when using data from the network or file system. The onus of trust is ultimately up to the developer and the network connections they make, but loading foreign data is inherently risky, and should not be used for input into sensitive operations. Developers are advised against the following:

Using data from a network source to determine a filename

Using data from a network source to construct a URL that the application uses to send private information

b) Risk from using, storing, or transmitting unsecured credentials

Storing user credentials on the user's local file system inherently introduces the risk that these credentials may be compromised. Developers are advised to consider the following:

If credentials must be stored locally, encrypt the credentials when writing to the local file system. Adobe AIR provides an encrypted storage unique to each installed application, which is described in detail in the developer documentation.

Do not transmit unencrypted user credentials to a network source unless that source is trusted.

Never specify a default password in credential creation — let users create their own. Users who leave the default expose their credentials to an attacker that already knows the default password.

c) Risk from a downgrade attack

During the application installation process, the runtime checks to ensure that a version of the application is not currently installed. If an application is already installed, the runtime compares the version string of the existing application against the version that is being installed. If this string is different, the user can choose to upgrade their installation.

The runtime cannot guarantee that the newly installed version is newer than the older version, only that it is different. An attacker can distribute an older version to the user to circumvent a security weakness. This risk can be mitigated, and the developer documentation provides best practices for implementing version schemes and update checks to avoid risks.